

# **Study Of Intel PML Effectiveness For Virtual Machine Working Set Size Estimation**

**PhD Student  
Stella Bitchebe**

**SupervisorS  
Pr. Alain Tchana  
Pr. Laurent Réveillère**

**Research team  
SCALE - I3S**

***Dresden-Germany, March'19***

# ● Introduction - Virtualization

- Basic building block in data centers


- Users reserve resources by booking VMs

- Overestimate VMs Sizes


A large percentage of VMs require many fewer resources than their allocations

(SOSP'17)

There is a significant potential for oversubscribing physical resources



Twitter clusters CPU utilization [SOSP'17]



Azur clusters Memory utilization [ASPLOS'14]

# On-demand allocation

## Context

- To avoid resource waste
- Needs accurate and efficient estimation
- Especially for memory: **limiting resource**

## Existing solutions

- All software base → Several drawbacks
- Examples
  - Self-Ballooning
  - Geiger
  - ZBalloon
  - **VMWare**

# On-demand allocation

## Context

- To avoid resource waste
- Needs accurate and efficient estimation
- Especially for memory: **limiting resource**

## Existing solutions

- All software base → Several drawbacks
- Examples
  - Self-Ballooning
  - Geiger
  - ZBalloon
  - **VMWare**

## VMWare Technique

- Periodically and randomly select  $n$  pages from the VM's memory and invalidate them
- Estimation is done by counting the number of pages which were subject to a non-present or read-only fault during the previous time interval

# On-demand allocation

## Context

- To avoid resource waste
- Needs accurate and efficient estimation
- Especially for memory: **limiting resource**

## Existing solutions

- All software base → Several drawbacks
- Examples
  - Self-Ballooning
  - Geiger
  - ZBalloon
  - **VMWare**

## VMWare Technique

- Periodically and randomly select  $n$  pages from the VM's memory and invalidate them
- Estimation is done by counting the number of pages which were subject to a non-present or read-only fault during the previous time interval





# INTEL PMIL



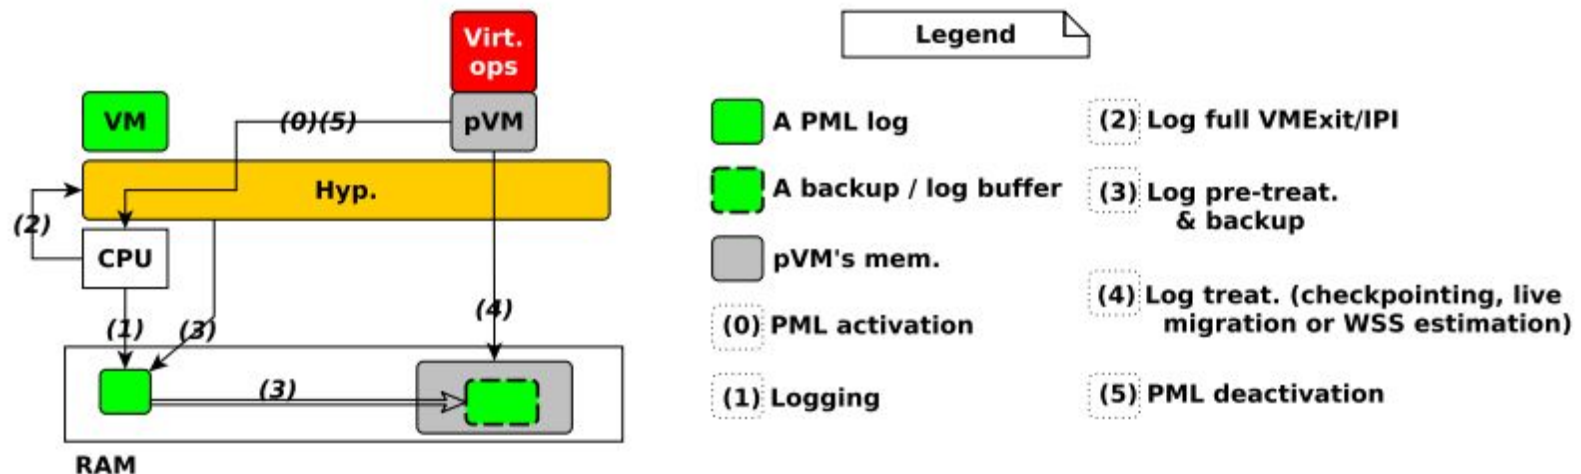
**Page Modification Logging**

“

*PML ... provides a faster mechanism for a hypervisor to monitor all the memory pages that a guest VM modifies, in order to improve working page set analysis, checkpointing, and even VM migration.*

*Richard Brunner, VMWare [[https://www.theregister.co.uk/2016/05/11/dirty\\_page\\_logs\\_coming\\_to\\_future\\_vsphere\\_release/](https://www.theregister.co.uk/2016/05/11/dirty_page_logs_coming_to_future_vsphere_release/)]*

# PML Architecture



- Once PML is enabled, if the VM modifies a page the CPU logs its GPA (Guest-Physical address) inside a buffer (*PML logging buffer*)
- When the log is full (512 addresses logged), the CPU raises a VMExit which traps inside the hypervisor
- The handler of the VMExit does certain task (e.g., copy the content of the *PML logging buffer* to a larger buffer which is shared with the pVM)
- Then the *PML index* is reset to 511 and the VM resumes



# Contribution

1

**Study of PML effectiveness for Live Migration, Checkpointing and Working Set Size estimation**

- **Xen and KVM Implemented PML**
- **Xen use PML for Live migration**

# Contribution

1

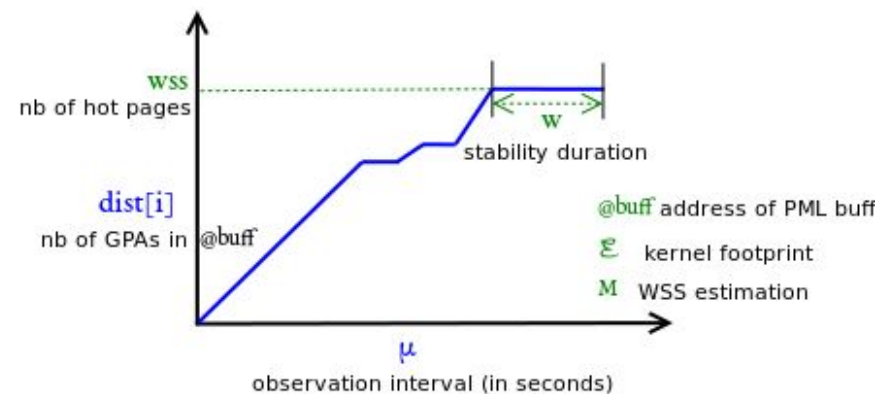
## Study of PML effectiveness for Live Migration, Checkpointing and Working Set Size estimation

- Xen and KVM Implemented PML
- Xen use PML for Live migration

2

## WSS estimation algorithm based on PML

- With Xen Hypervisor



$$M = wss * size\_of\_a\_page + \epsilon$$



## Testbed

On this machine:  
PML is not yet  
present on servers



## Testbed

On this machine:  
PML is not yet  
present on servers

VM with single vCPU  
and 5GB (for micro  
bench)/10 (for macro  
benchs) memory



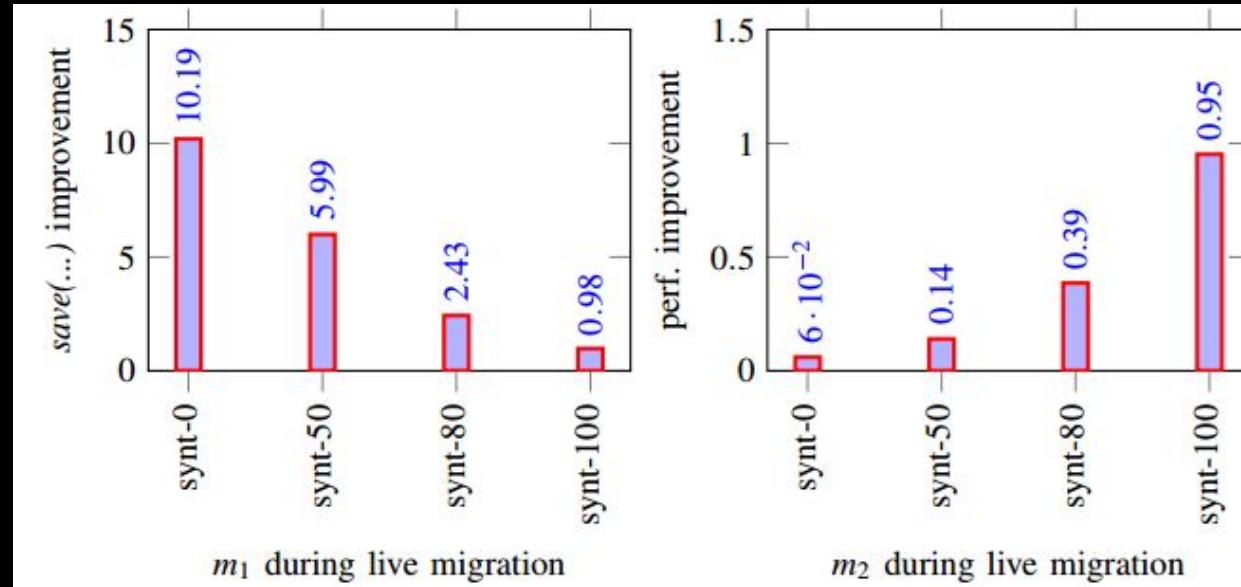
## Testbed

On this machine:  
PML is not yet  
present on servers

VM with single vCPU  
and 5GB (for micro  
bench)/10 (for macro  
benches) memory

Micro Bench: A  
synthetic application  
written for the  
purpose

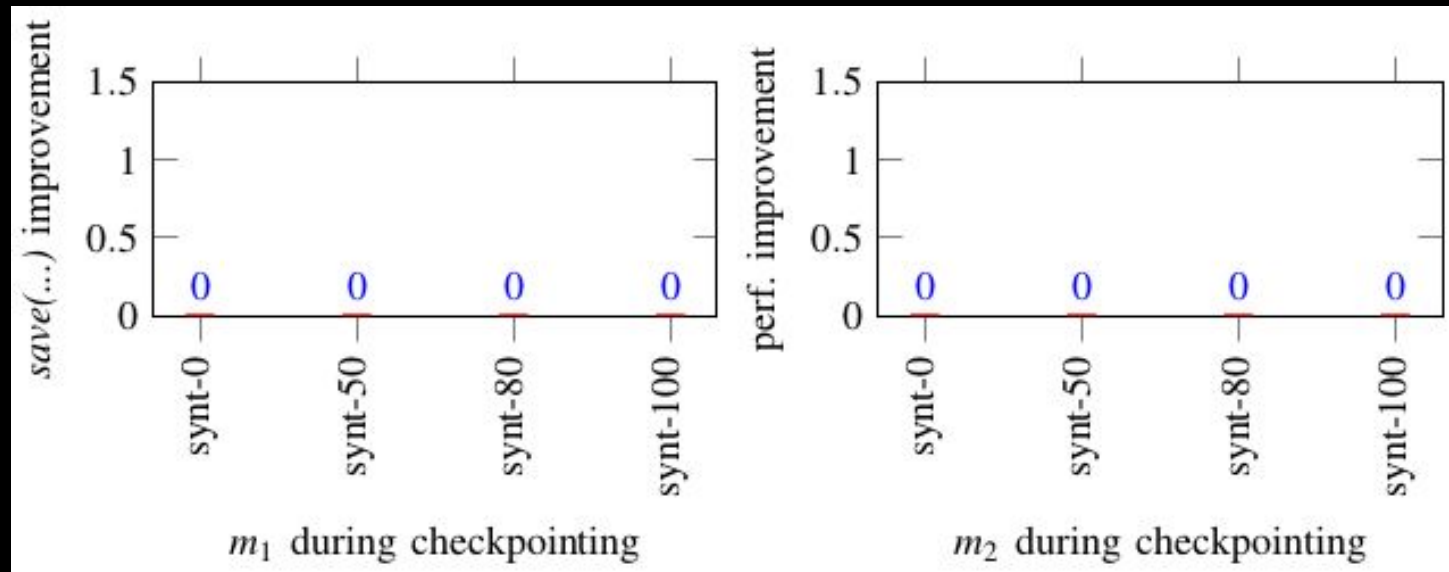
## Results for Live migration



- **Baseline:** classical approach (write protecting)
- **synt-x:** x is the write intensity (100 means write only and 0 means read only)
- Average of 10 executions
- Metric 1 ( $m_1$ ): impact on live migration duration
- Metric 2 ( $m_2$ ): impact on user application performance

- ➔ PML reduces the duration of live migration by 0.98%-10.18%: especially for read intensive applications
- ➔ PML reduces the impact of live migration on user applications by 0.065%-0.95%: especially for write intensive applications

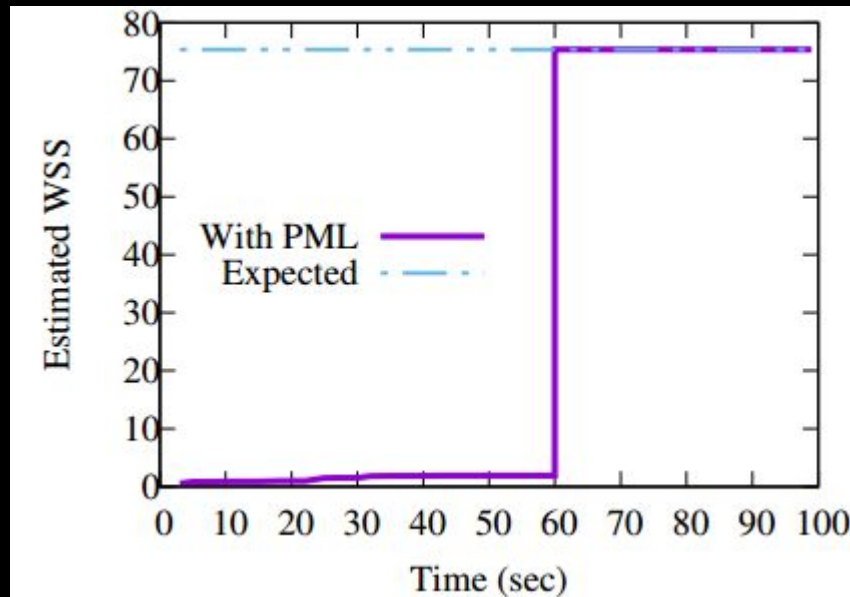
## Results for Checkpointing



→ PML does not improve checkpointing

→ Live Checkpointing is likely to take advantage of PML

## Results for WSS estimation




**PML Hard Limit - PML is not able to accurately estimate a VM WSS**

- ➔ PML does not log accessed pages
- ➔ With PML it is not possible to track hot pages




## Results for WSS estimation



**PML Hard Limit - PML is not able to accurately estimate a VM WSS**

- ➔ PML does not log accessed pages
- ➔ With PML it is not possible to track hot pages



**PML Soft Limit - PML is unfair for cloud users and their VMs**

- ➔ The VMExit on *PML log buffer full* is handled by the CPU that runs the VM for which WSS is estimated



## Conclusion - Ongoing Work

- Proposed Extension of PML: PRL (Page Reference Logged)
- PRL implementation under Gem5 simulator
- PRL-based WSS estimation algorithm

That was all!

- **ANY QUESTIONS?**