

Out of Hypervisor (OoH): When Nested Virtualization Becomes Practical

Stella Bitchebe

PhD Thesis Defense
February 3rd, 2023



Reviewers

Pr Edouard Bugnion, EPFL
Pr Gaël Thomas, Telecom SudParis

Examiners

Pr Guillaume Urvoy-Keller, UCA

Pr Laurent Réveillère, Univ. Bordeaux

Dr Natacha Crooks, UC Berkeley

Dr Oana Balmau, McGill Univ.

Dr Renaud Lachaize, Univ. Grenoble

Invited

Dr Anne-Marie Chana, ENSPY

Dr Bernabe Batchakui, ENSPY

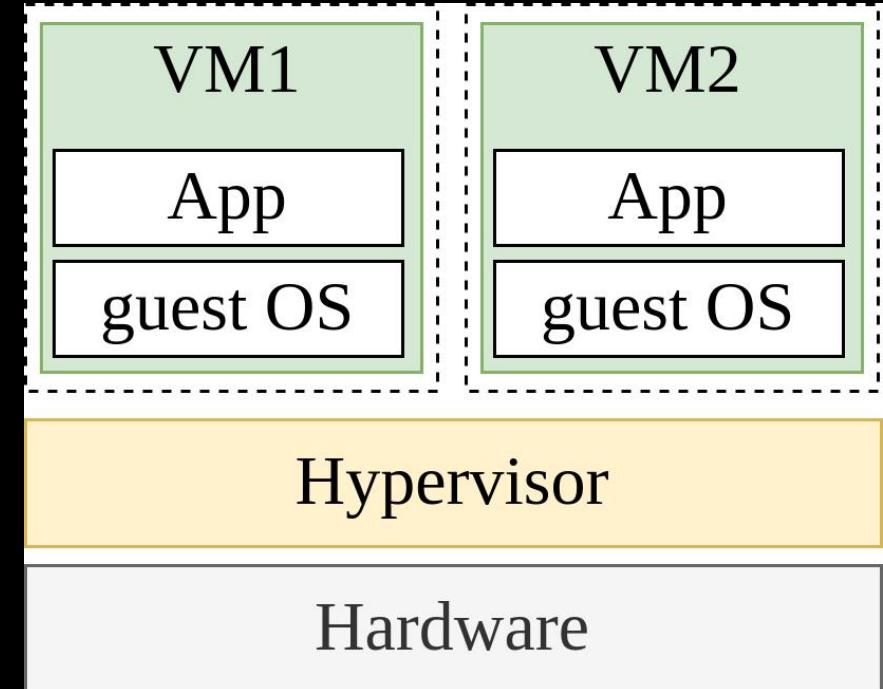
Advisor

Pr Alain Tchana, Univ. Grenoble

Nested Virtualization



Definition



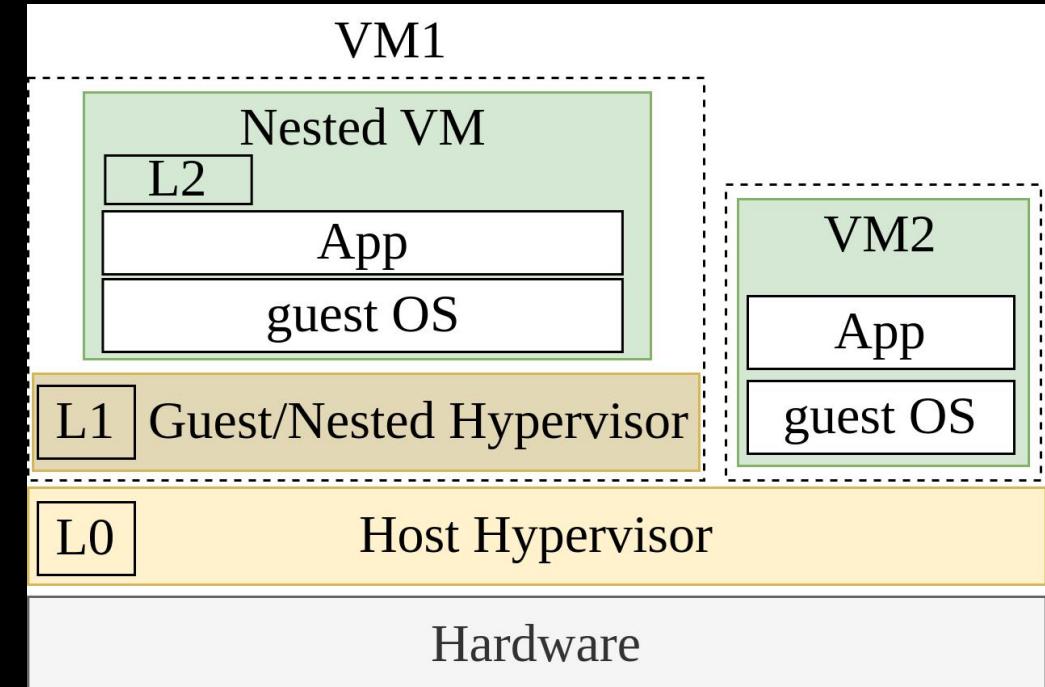
Nested Virtualization



Definition

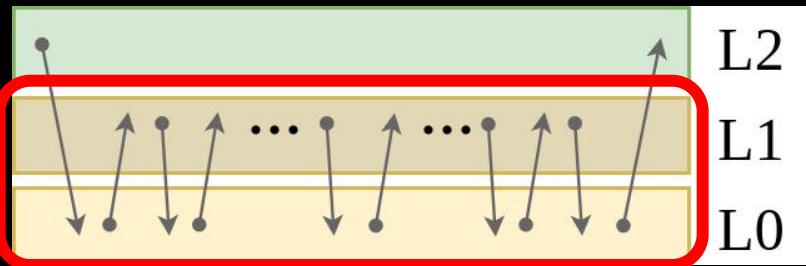
“Recursive virtualization, where a virtual machine runs under itself a copy of a hypervisor.”

Popek and Goldberg, 1970s



Nested Virtualization

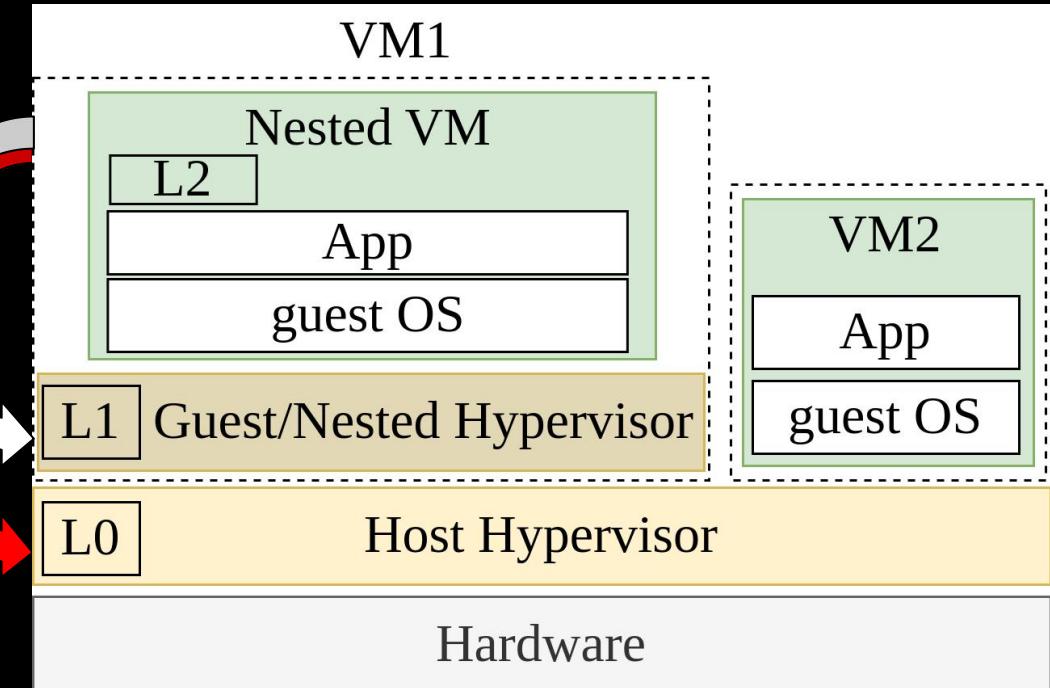
⚠ Drawbacks



non-nested exit

nested exit

Single trap from nested VM produces multiple traps
to the host hypervisor



Significant number of VM exits (at least $\times 2$)

Vilanova et al. [150, ISCA'19]

Nested Virtualization



Use Cases of Nested Virt.



Tests/Development/Demo

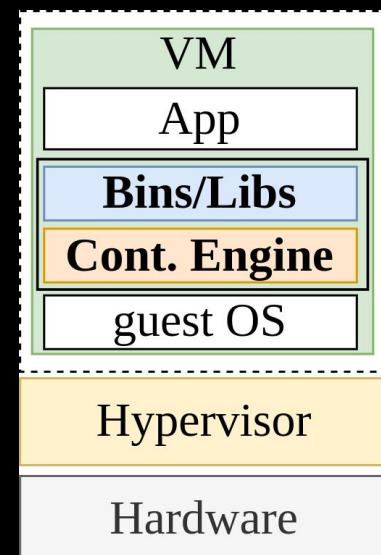
- No matter the overhead
- Not in production



Isolation



Containers



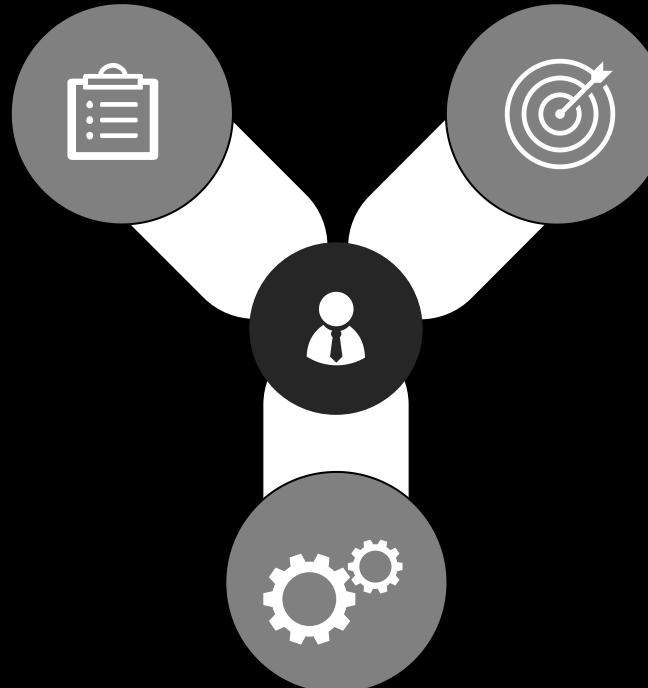
OoH: Out of Hypervisor



Presentation

What?

New virtualization
research axis



Why?

Avoid trying to virtualize full
virtual hardware inside a VM

How?

Expose individually hypervisor oriented hardware virt. features to the guest OS so that its processes can benefit from them

OoH: Out of Hypervisor



Virtualization Technologies Categorization

G1: Resource Multiplexing

- ✓ EPT (*Extended Page Table*): MMU virtualization
- ✓ SRIOV (*Single Root I/O Virtualization*)
- ✓ APICv (*Advanced Programmable Interrupt Controller virtualization*)
- ✓ etc.

G2: VM's Management

- ✓ PML (*Page Modification Logging*)
- ✓ CAT (*Cache Allocation Technology*)
- ✓ SPP (*Sub-Page write Permissions*)
- ✓ etc.

OoH: Out of Hypervisor



Virtualization Technologies Categorization

G1: Resource Multiplexing

- ✓ EPT (*Extended Page Table*): MMU virtualization
- ✓ SRIOV (*Single Root I/O Virtualization*)
- ✓ APICv (*Advanced Programmable Interrupt Controller virtualization*)
- ✓ etc.

G2: VM's Management

- ✓ PML (*Page Modification Logging*)
- ✓ CAT (*Cache Allocation Technology*)
- ✓ SPP (*Sub-Page write Permissions*)
- ✓ etc.

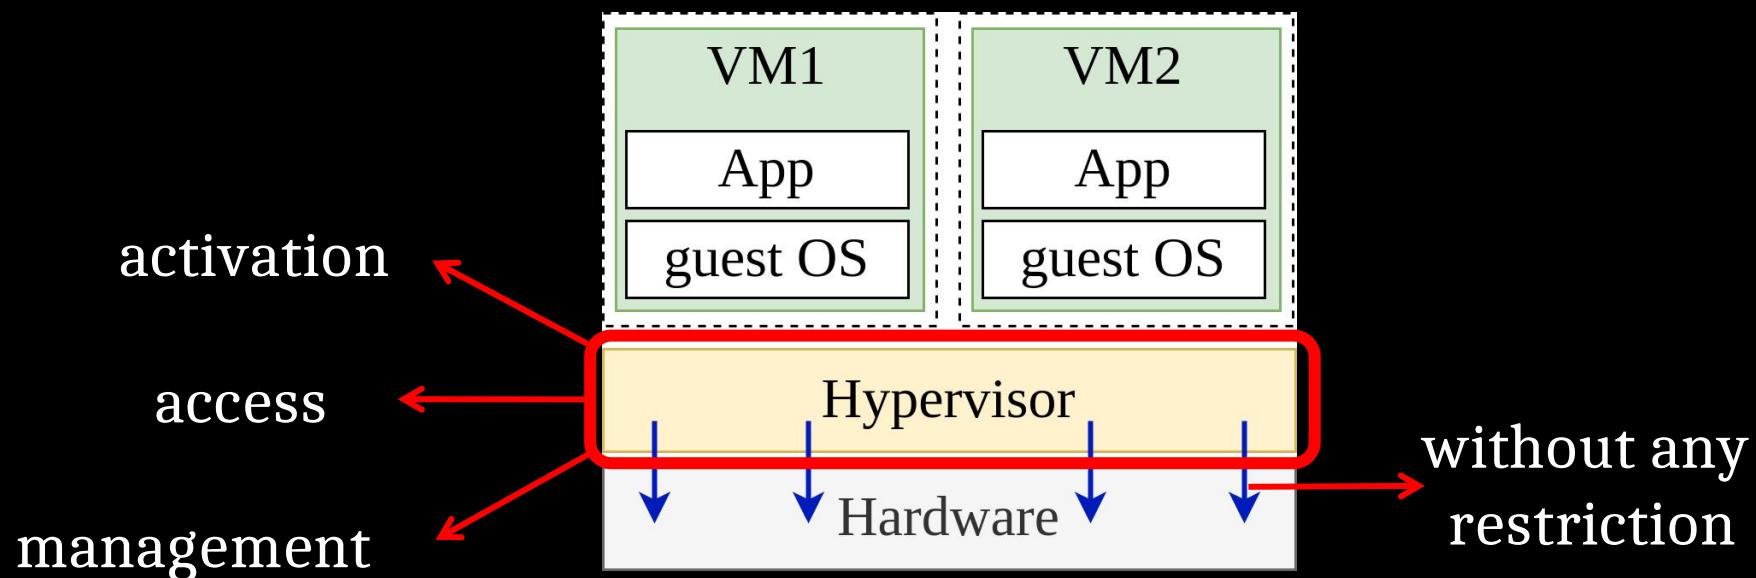


OoH scope: Can be exploited to remove the need for a nested hypervisor to leverage hypervisor properties

Problem Statement



All hardware functionalities are hypervisor-oriented



OoH's State-of-the-art



State-of-the-art

Leveraging HAV: Dune

Direct Access to HW: Hardened Hypervisors

Improving Nested Virt.

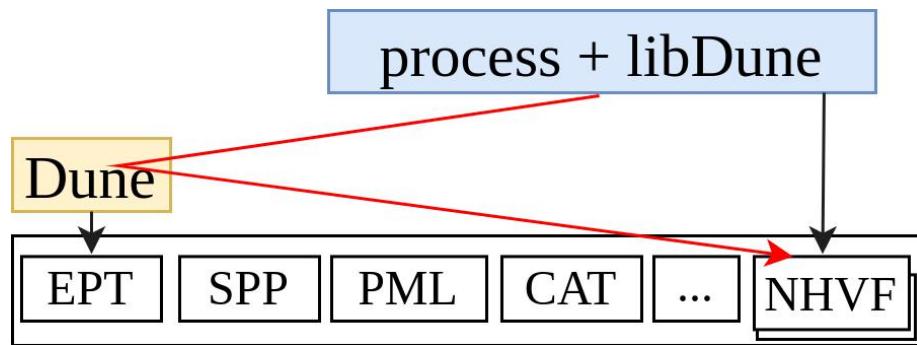
OoH's State-of-the-art



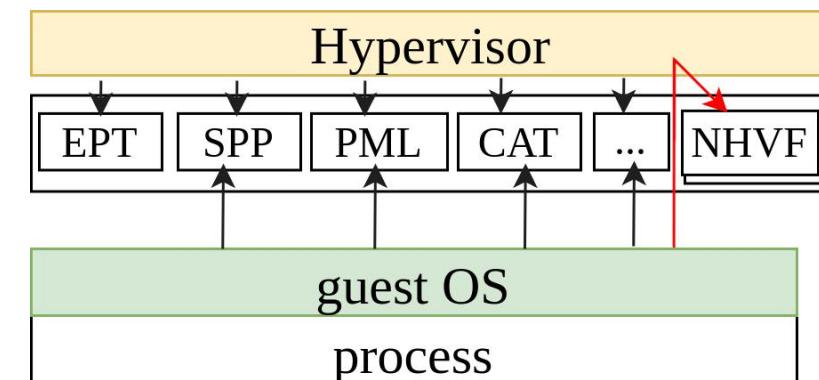
State-of-the-art

Leveraging HAV: Dune (OSDI 2012)

Dune



OoH



Dune uses hardware virt. features (EPT) to allow processes in bare-metal to access non-virt. feat. (execute privileged instructions) while OoH makes processes in VMs to access and use hardware virt. features

OoH's State-of-the-art



State-of-the-art

Direct Access to HW: Hardened Hypervisors

Nohype (ISCA 2010):

1CPU core per VM, fixed memory per VM, etc.

AWS Nitro (Amazon 2017):

network, storage, and security management from the hypv. to the HW (-> Nitro cards) to enable near-native perf



Restricted #instances

Rigid hypervisor

OoH's State-of-the-art



State-of-the-art

Improving Nested Virt.

NEVE (Nested Virtualization Extensions for ARM, SOSP 2017, Lim et al.): access to VM data structures to reduce exits multiplication from nested hypv. → limited fields accessible

SV_T (SMT-based VirTualization, ISCA 2019, Vilonava et al.): pins each virt. layer to a single core → limits resources sharing

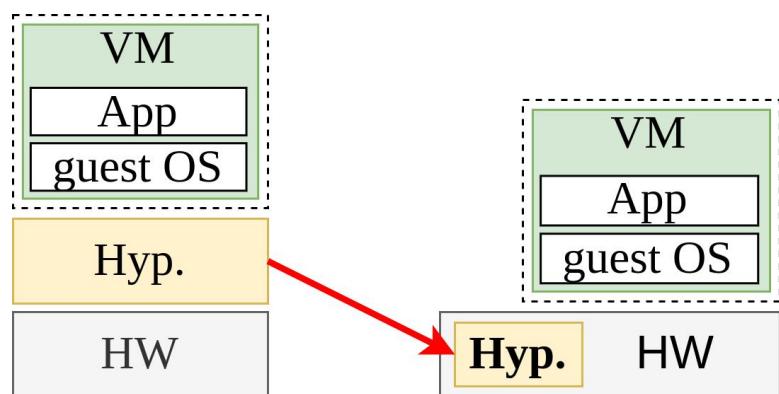
DVH (Direct Virtual Hardware, ASPLOS 2020, Lim et al.): direct IO devices and timers to VMs → unpractical application to all devices

OoH's State-of-the-art



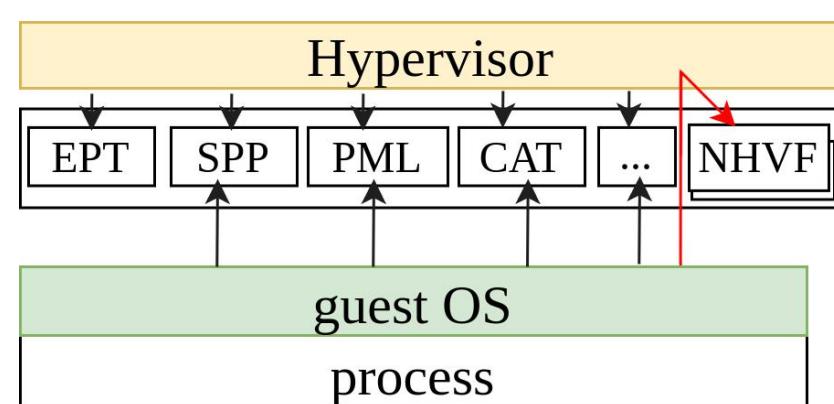
OoH in the virtualization landscape

Hardened Hyp.



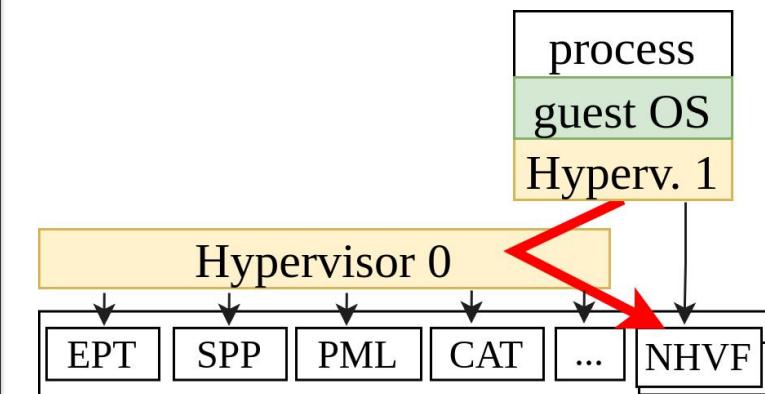
- ✓ No nested hyp.
- ✗ Rigid hypervisor
- ✗ Overcommitment unpractical

OoH



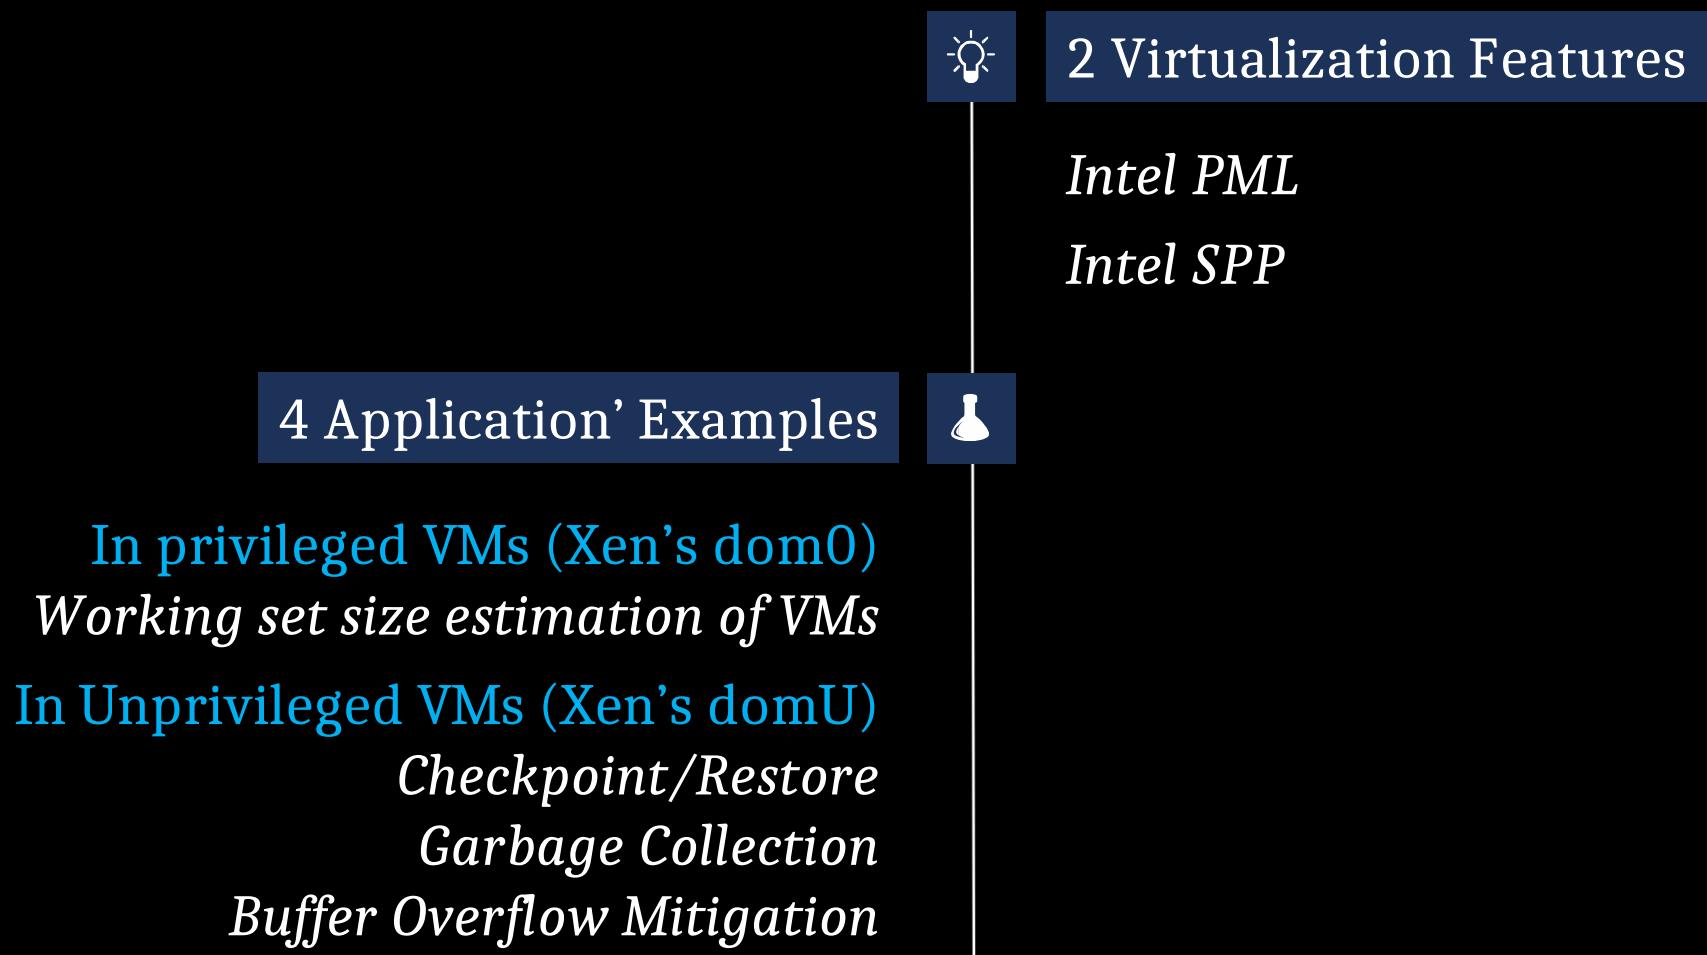
- ✓ No nested hyp.
- ✓ Flexible light kernel module
- ✓ Overcommitment possible

Nested virt.



- ✗ Kept nested hyp.
- ✓ Overcommitment possible

💡 OoH illustration in this thesis



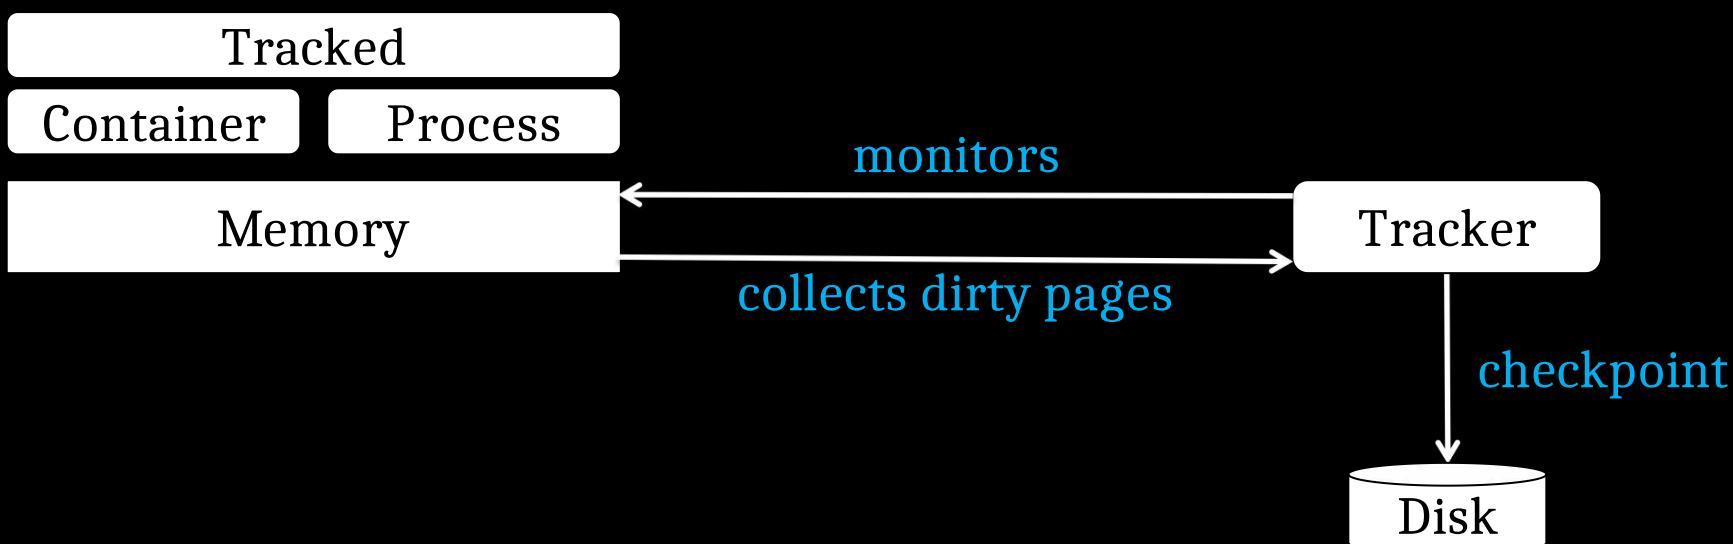
OoH for PML: Efficient Dirty Page Tracking In Virtualized Clouds

Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking

PML is exposed to the guest for HW-assisted **dirty page tracking**: monitoring of dirty (or modified) pages of an application/process

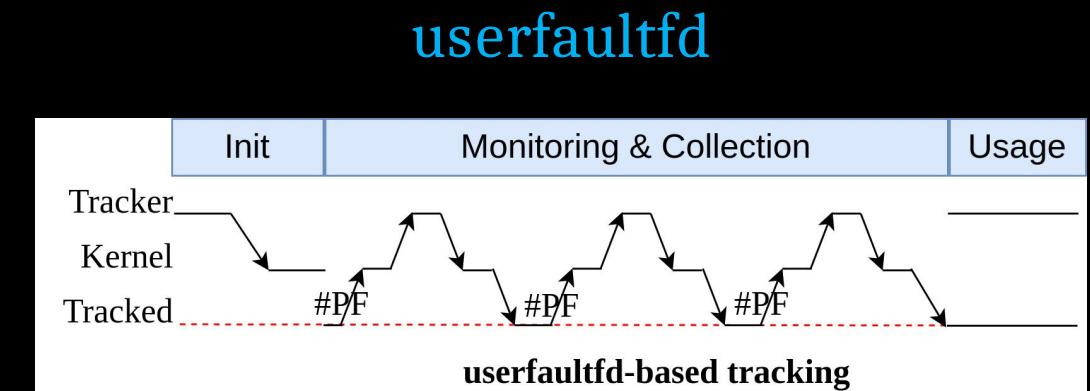
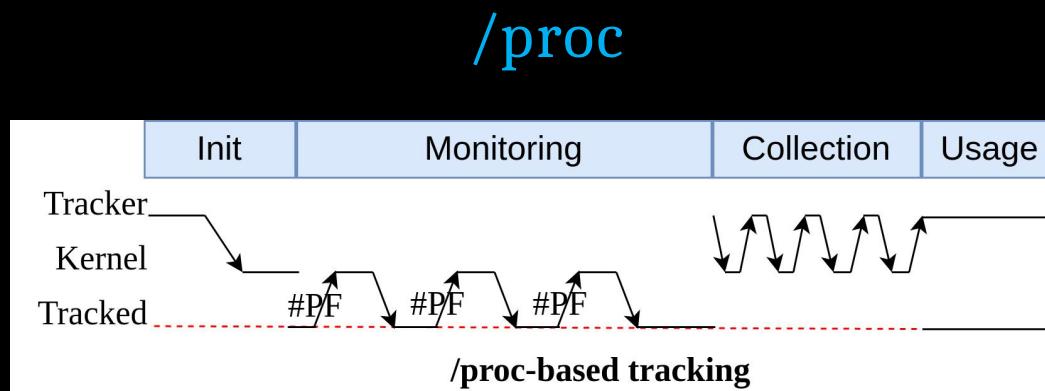
Usefulness of dirty page tracking: checkpointing (backup images for recovery after failure), garbage collection (freeing unused objects for better memory management)



Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking: Current Approach

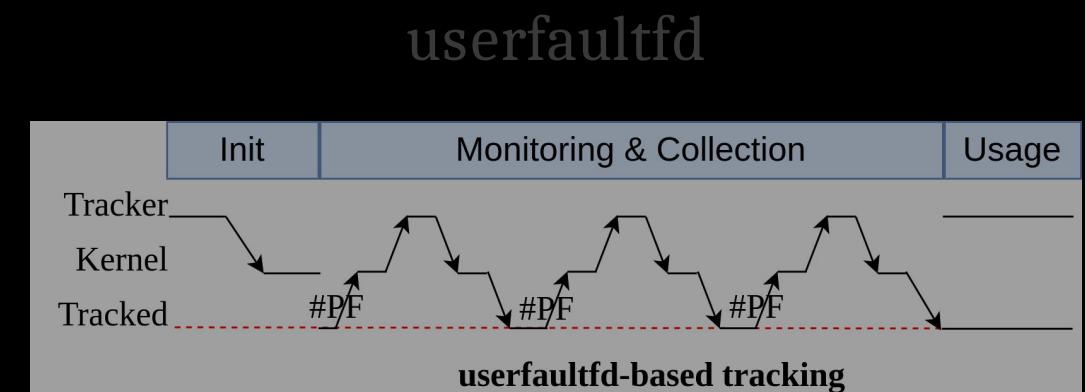
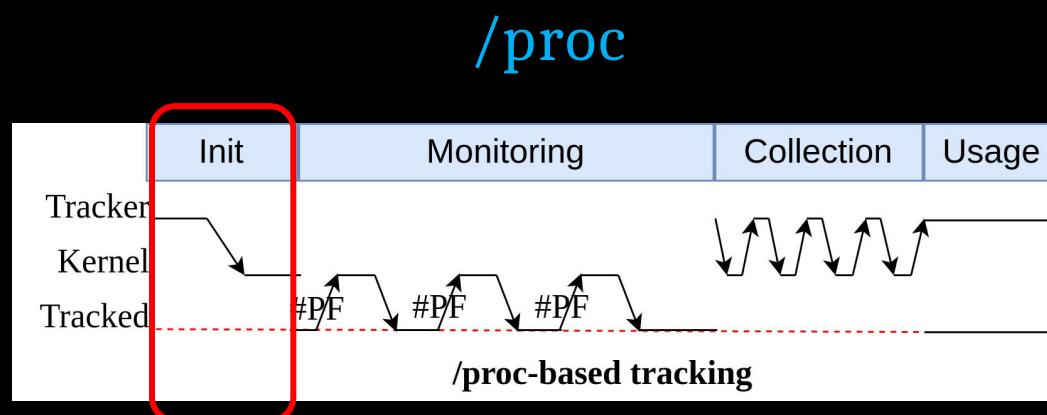
Page write-protection
2 main interfaces in Linux



Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking: Current Approach

Page write-protection
2 main interfaces in Linux



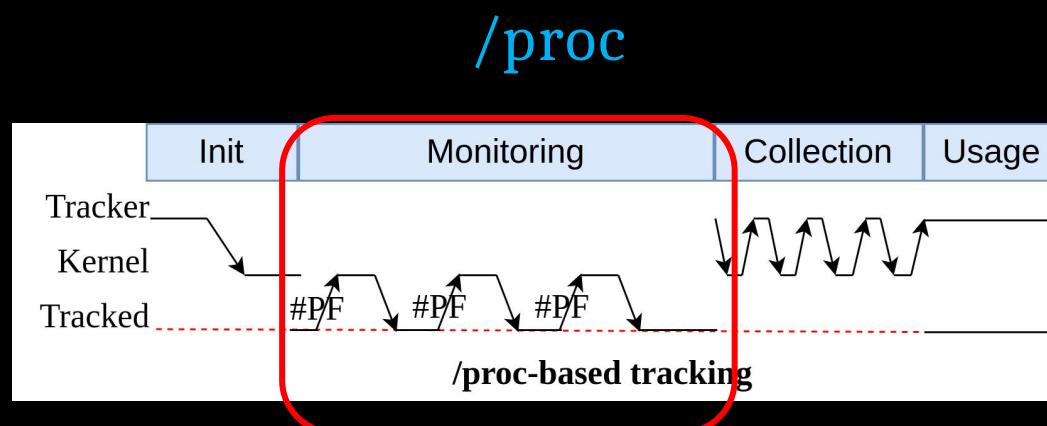
syscall to invalidate all 'Tracked' pages

Efficient Dirty Page Tracking In Virtualized Clouds

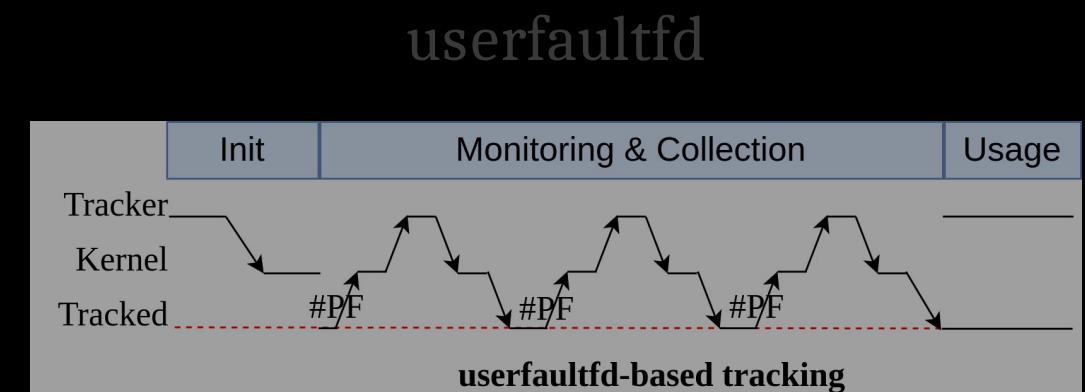
📍 Dirty Page Tracking: Current Approach

Page write-protection

2 main interfaces in Linux



page faults (#PF) generated upon access

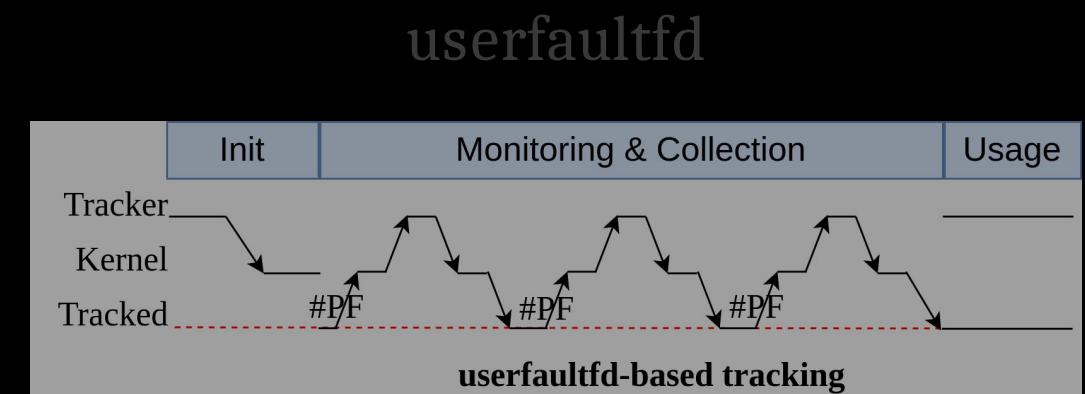
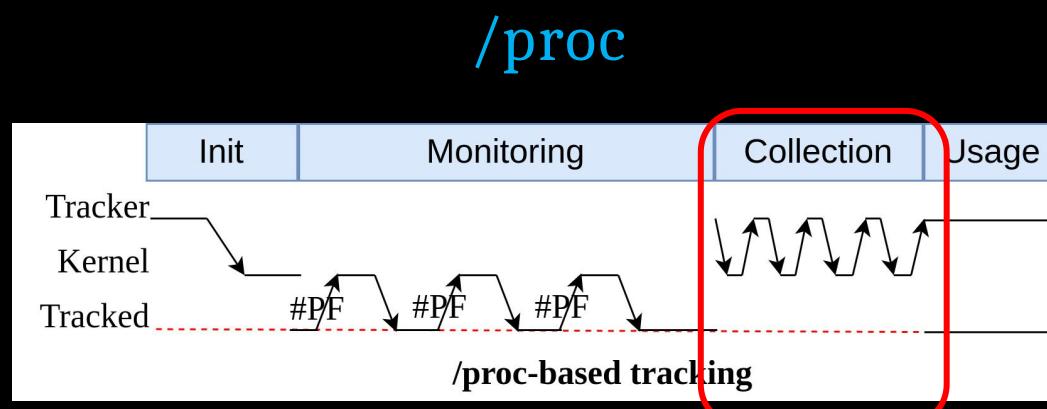


Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking: Current Approach

Page write-protection

2 main interfaces in Linux

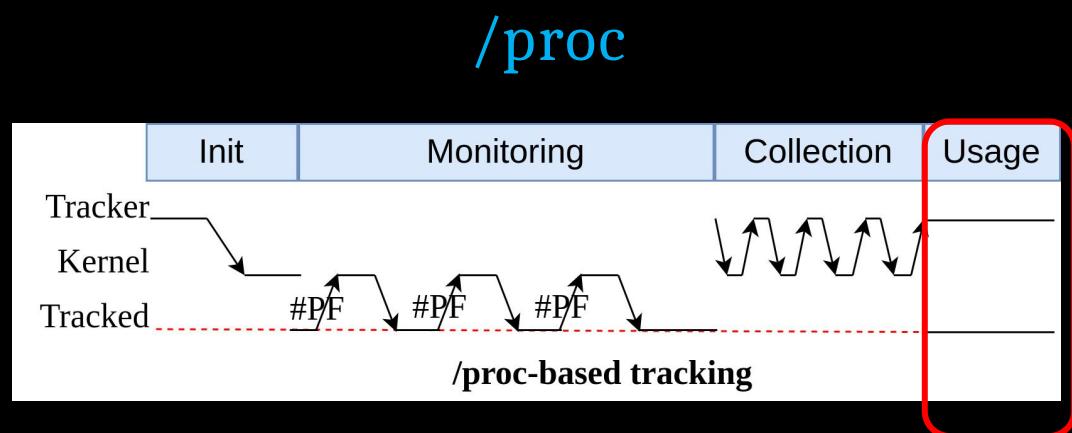


Tracker reads `/proc/PID/maps` and `/proc/PID/pagemap` (with PID of Tracked) uses the dirty bit to determine accessed pages

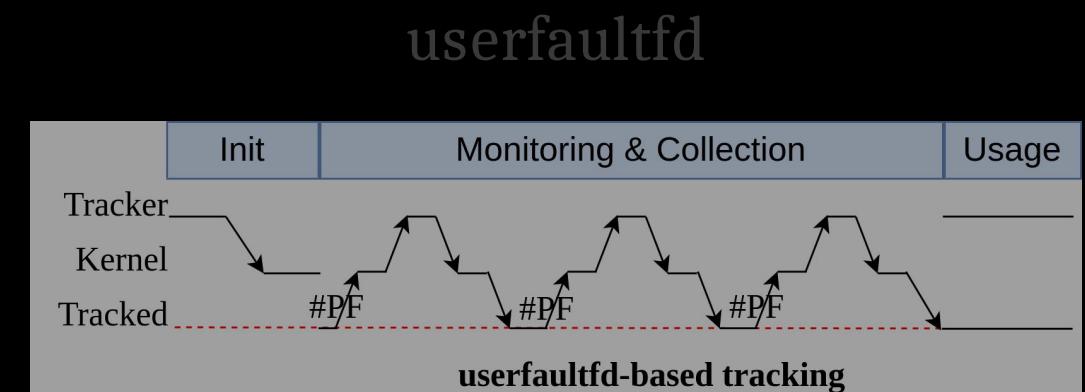
Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking: Current Approach

Page write-protection
2 main interfaces in Linux



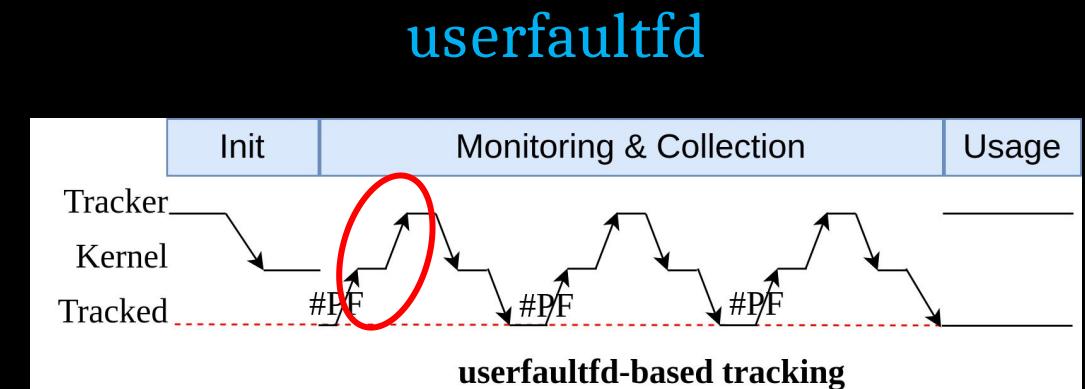
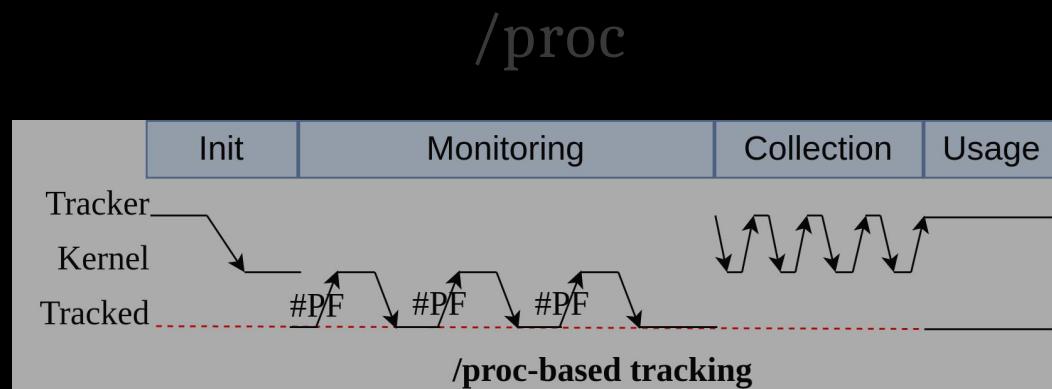
e.g., for checkpointing



Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking: Current Approach

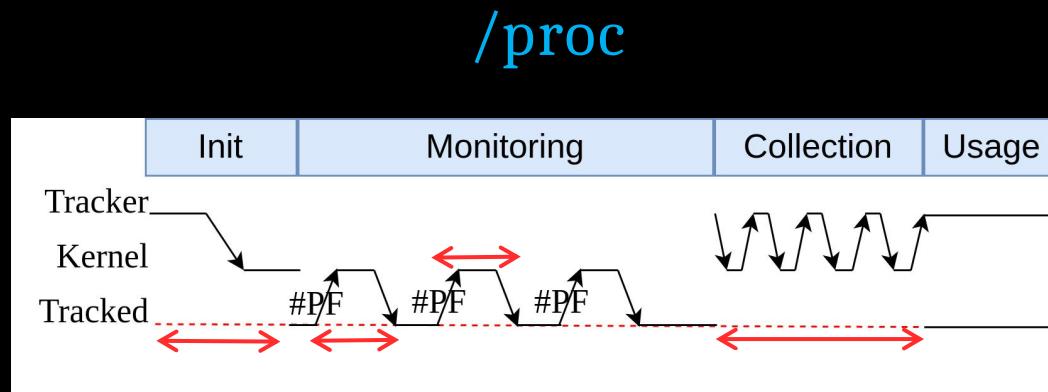
Page write-protection
2 main interfaces in Linux



the kernel transfers the #PF handling to userspace

Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking: Current Approach

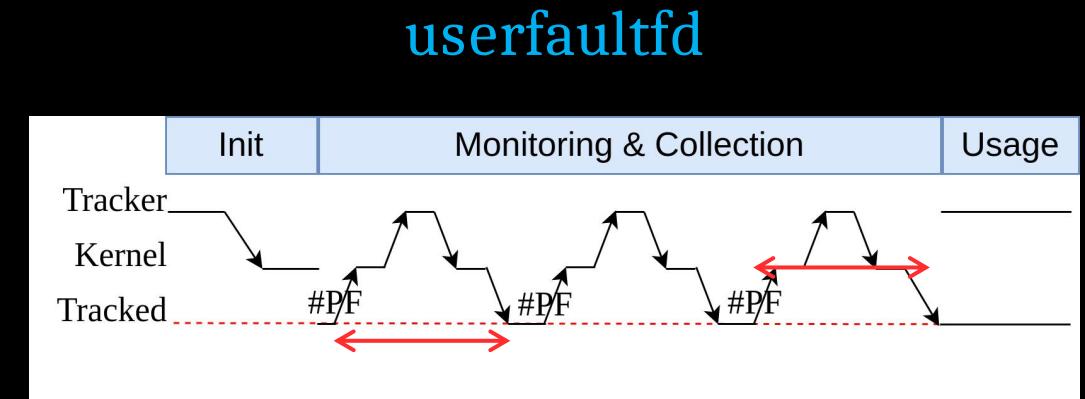


For 1GB working set:

PT walk:
in kernel (with TLB flush): ~2.23ms
in userspace (/proc/PID/pagemap): ~594.18ms



4x slowdown on Tracked
2.5x slowdown on Tracker



For 1GB working set:

#PF handling and context switches

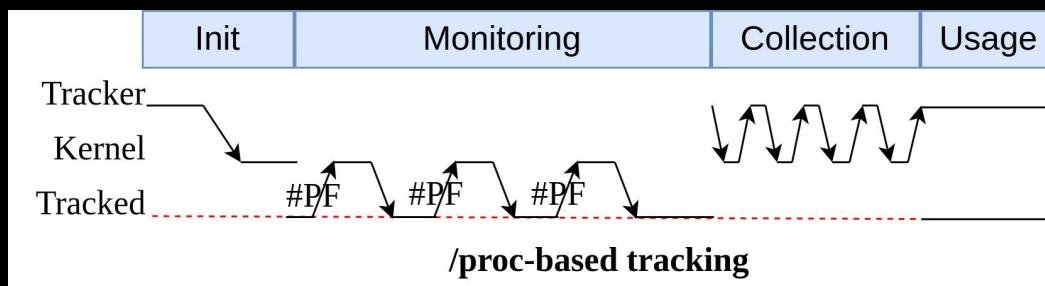


15.6x slowdown on Tracked
14.5x slowdown on Tracker

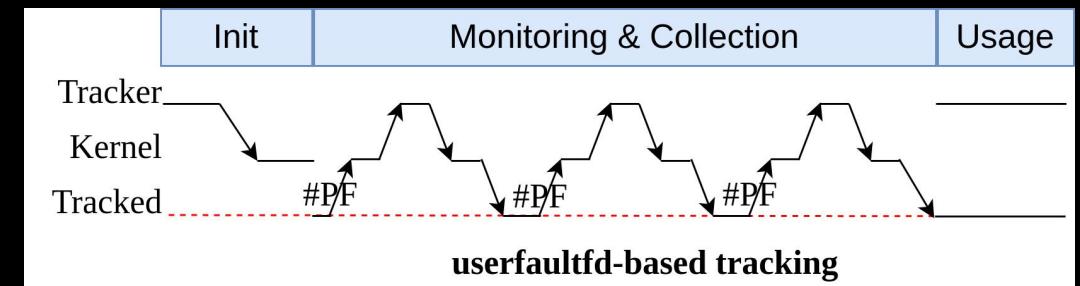
Efficient Dirty Page Tracking In Virtualized Clouds

📍 Dirty Page Tracking in Virt. Clouds: PML-based

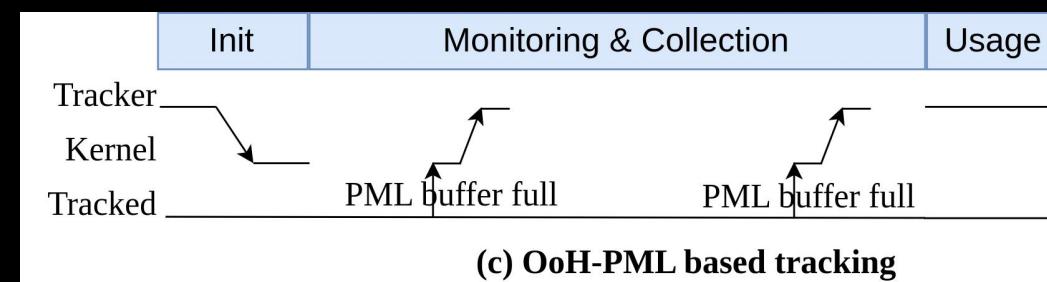
/proc



userfaultfd



PML



Efficient Dirty Page Tracking In Virtualized Clouds

Challenges

(C₁) Only the hypervisor can manage PML

Efficient Dirty Page Tracking In Virtualized Clouds

Challenges

(C₁) Only the hypervisor can manage PML

(C₂) PML works at coarse-grained => it concerns the entire VM

Efficient Dirty Page Tracking In Virtualized Clouds

Challenges

(C₁) Only the hypervisor can manage PML

(C₂) PML works at coarse-grained => it concerns the entire VM

(C₃) PML only logs GPAs

Efficient Dirty Page Tracking In Virtualized Clouds

Solutions

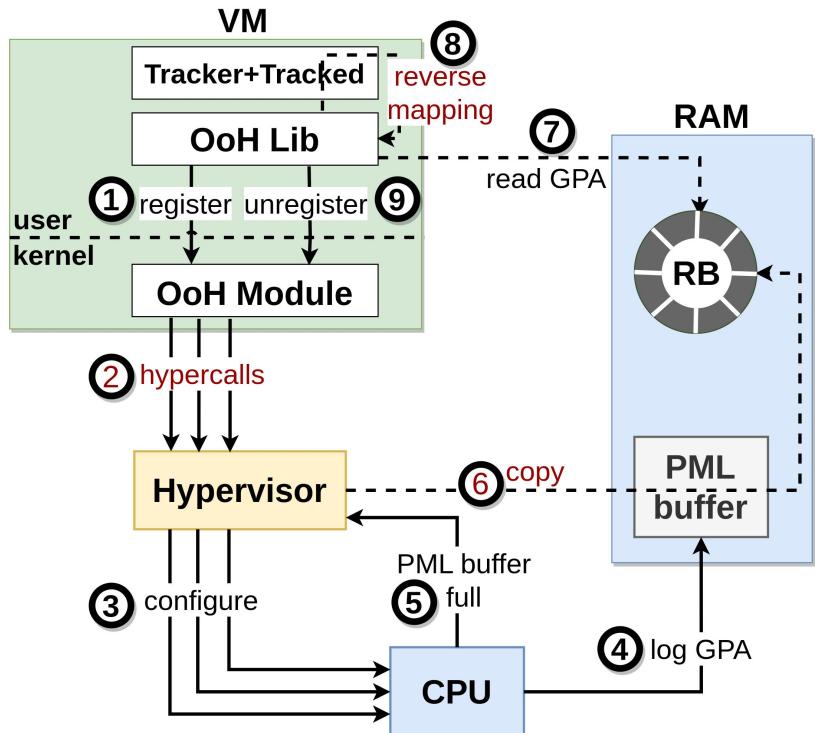
Shadow PML (SPML)

Extended PML (EPML)

Efficient Dirty Page Tracking In Virtualized Clouds

Solutions

Shadow PML (SPML)



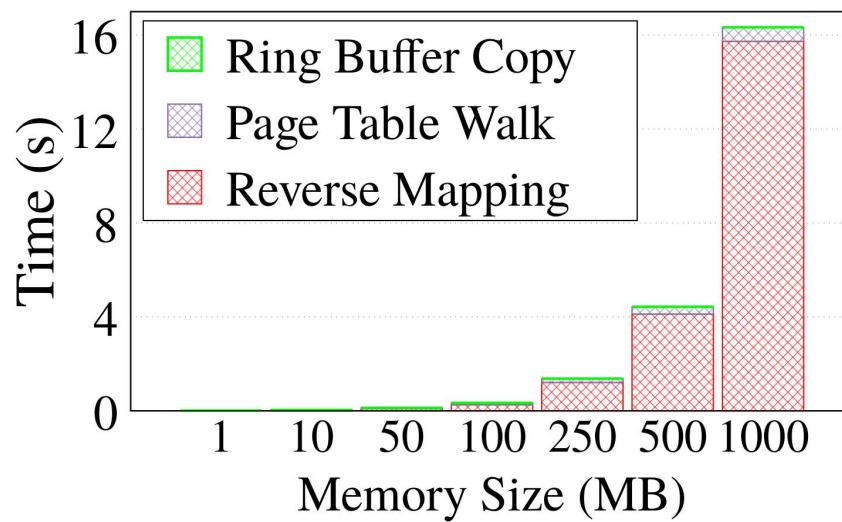
Extended PML (EPML)

Efficient Dirty Page Tracking In Virtualized Clouds

Solutions

Shadow PML (SPML)

Costly reverse mapping (~15.739 s for 1GB working set)



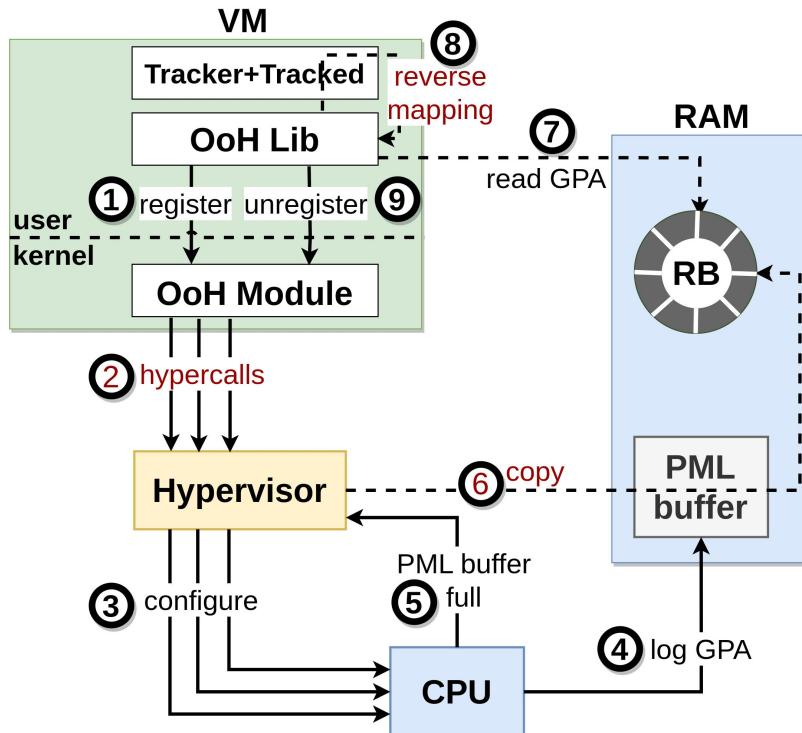
Extended PML (EPML)

Costly hypercalls (4.49 μ s for an empty hypercall)

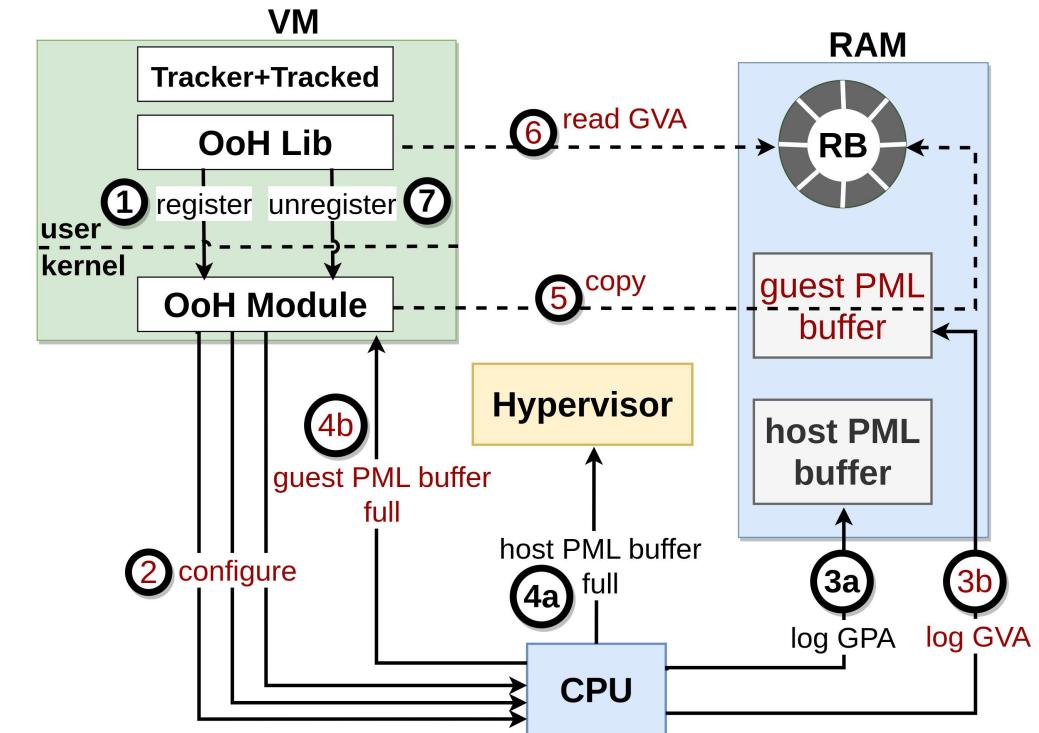
Efficient Dirty Page Tracking In Virtualized Clouds

Solutions

Shadow PML (SPML)



Extended PML (EPML)



Efficient Dirty Page Tracking In Virtualized Clouds

📍 Isolation & Security

Vis-à-vis the Hypervisor

- ✓ Small TCB 1 (194LOC) - at least safe as existing hypercalls
- ✓ Guest does not see nor manipulate host physical memory
- ✓ Ring buffer allocated from VM's memory

Efficient Dirty Page Tracking In Virtualized Clouds

📍 Isolation & Security

Vis-à-vis the Hypervisor

- ✓ Small TCB 1 (194LOC) - at least safe as existing hypercalls
- ✓ Guest does not see nor manipulate host physical memory
- ✓ Ring buffer allocated from VM's memory

Between VMs

- ✓ Intrinsic VM' isolation not altered
- ✓ Ring buffer allocated per VM's address space => no possible inference
- ✓ Per process ring buffer and restriction to tracker process only

Efficient Dirty Page Tracking In Virtualized Clouds

📍 Implementation

EPML's hardware changes under BOCHS

- ✓ EPML: 44 LOCs in 6 files

Xen as the hypervisor

- ✓ SPML: 182 LOCs in 13 files
- ✓ EPML: 120 LOCs in 9files

Linux as the guest OS

- ✓ SPML: 6 LOCs in 2 files
- ✓ EPML: 14 LOCs in 9files

OoH Module:

- ✓ SPML: +520 LOCs
- ✓ EPML: +520 LOCs

Integration of the OoH Lib with:

- ✓ CRIU: Checkpoint/Restore in User space (used in OpenVZ, Docker, etc.)
 - Based on /proc technique
 - SPML: 251 LOCs in 4 files
 - EPML: 140 LOCs in 4 files
- ✓ Boehm GC: popular C/C++ garbage collector (used in Mozilla, GNU Java Compiler, etc.)
 - Based on /proc technique
 - SPML: 254 LOCs in 4 files
 - EPML: 144 LOCs in 4 files

Efficient Dirty Page Tracking In Virtualized Clouds



Evaluation Methodology

Build a formula to evaluate EPML

Show its accuracy on other techniques that are measurable

Impact on Tracker: Execution time of Tracker when implementing technique x

$$E(C_{tker}) = E(C_x) + E(C_p) + I(C_x, C_p)$$

$x : /proc, SPML, EPML - C x : enable_PML, ring buffer copy, etc.$

Impact on Tracked: Time of Tracked when monitored by a Tracker using technique x

$$E(C_{tked_tker}) = E(C_{tked}) + E(C_{tker}) + I(C_x, C_{tked})$$

$I(C_x, C_{tked}): page faults, vmexits, etc., etc.$

Efficient Dirty Page Tracking In Virtualized Clouds

Evaluation Methodology: Validation

SPML

Metric	Time (ms)
$E(C_{t_{ker}})$ measured	5503.79
$E(C_{tked_t_{ker}})$ measured	135255.35
$E(C_p)$	251.35
$E(C_{copy_rb})$	0.49
$E(C_{disable\ pml})$	2.06
$E(C_{rev.\ mapping})$	5419
$E(C_{t_{ker}})$ estimated	5672.9
$E(C_{vmexits})$	18000
N	39
$E(C_{vmread,vmwrite})$	1.73×10^{-3}
$E(C_{tked_t_{ker}})$ estimated	136919.85

/proc

Metric	Time (ms)
$E(C_{t_{ker}})$ measured	1097.99
$E(C_{tked_t_{ker}})$ measured	115283.35
$E(C_p)$	251.35
$E(C_{clear_refs})$	1.409
$E(C_{PT\ walk})$	0.89
$E(C_{t_{ker}})$ estimated	1116.09
$E(C_{PFHkernel})$	0.27
$E(C_{tked_t_{ker}})$ estimated	114418.58

Efficient Dirty Page Tracking In Virtualized Clouds

Evaluation Methodology: Validation

SPML

Metric	Time (ms)
$E(C_{t_{ker}})$ measured	5503.79
$E(C_{tked_t_{ker}})$ measured	135255.35
$E(C_p)$	251.35
$E(C_{py_rb})$	0.49
$E(C_{\text{enable pml}})$	2.06
$E(C_{\text{rev. mapping}})$	5419
$E(C_{t_{ker}})$ estimated	5672.9
$E(C_{vmexits})$	18000
N	39
$E(C_{\text{vmread,vmwrite}})$	1.73×10^{-3}
$E(C_{tked_t_{ker}})$ estimated	136919.85

96.34%

accuracy

Metric	Time (ms) /proc
$E(C_{t_{ker}})$ measured	1097.99
$E(C_{tked_t_{ker}})$ measured	115283.35
$E(C_p)$	251.35
$E(C_{clear_refs})$	1.409
$E(C_{PT \text{ walk}})$	0.89
$E(C_{t_{ker}})$ estimated	1116.09
$E(C_{\text{PFHkernel}})$	0.27
$E(C_{tked_t_{ker}})$ estimated	114418.58

99%
accuracy

Efficient Dirty Page Tracking In Virtualized Clouds



Evaluation Benchmarks

Benchmarks

Macro-benchmarks: tkrzw applications (key value store) and Phoenix applications MapReduce)

Three working set sizes (Small, Medium, and Large)

Metrics

Execution times of tracker and tracked

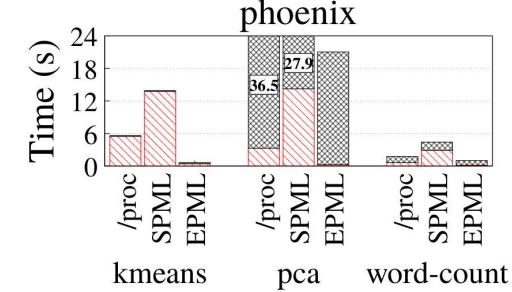
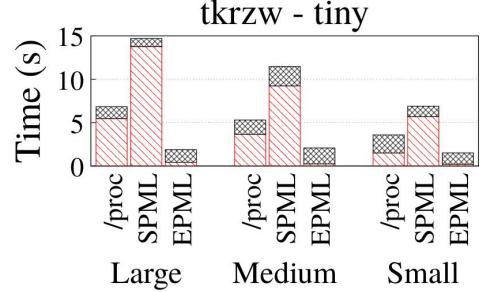
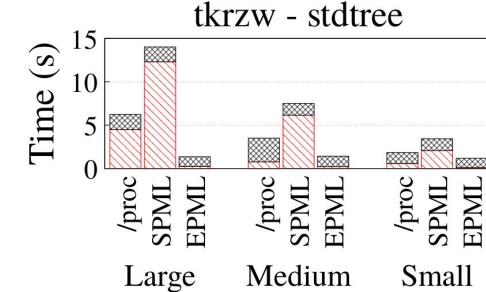
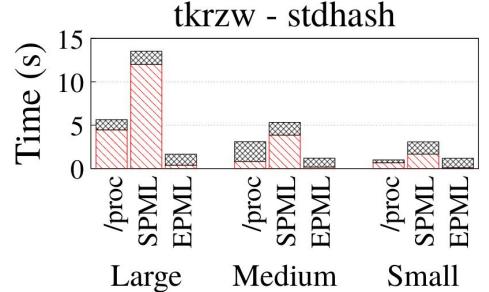
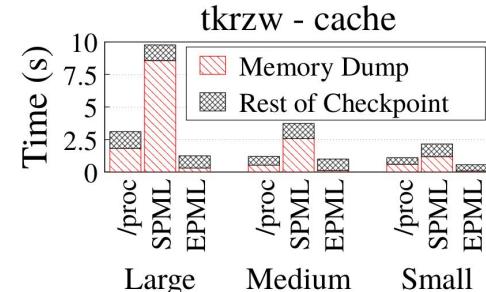
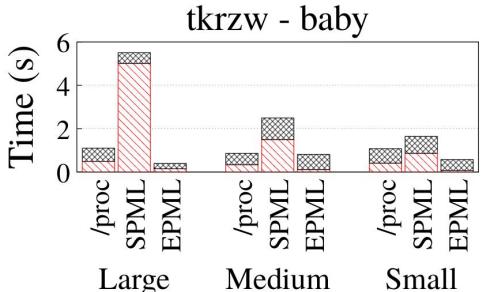
Evaluations aim answering the following questions:

1. What is the potential overhead or improvement of SPML and EPML compared to existing solutions (/proc and ufd)?
2. What is the scalability of SPML and EPML?

Efficient Dirty Page Tracking In Virtualized Clouds



Evaluation Results with CRIU: for Tracker



Mem dump (MD) phase: during which CRIU collects dirty pages

- /proc => walk /proc/\$pid/(maps, pagemap)
- SPML => copy ring buffer + reverse mapping
- EPML => copy via ring buffer

SPML:

- reverse mapping is **+66%** of MD time
- **5X** slowdown vs. /proc for tkrzw-baby and phoenix-kmeans (config. Large)

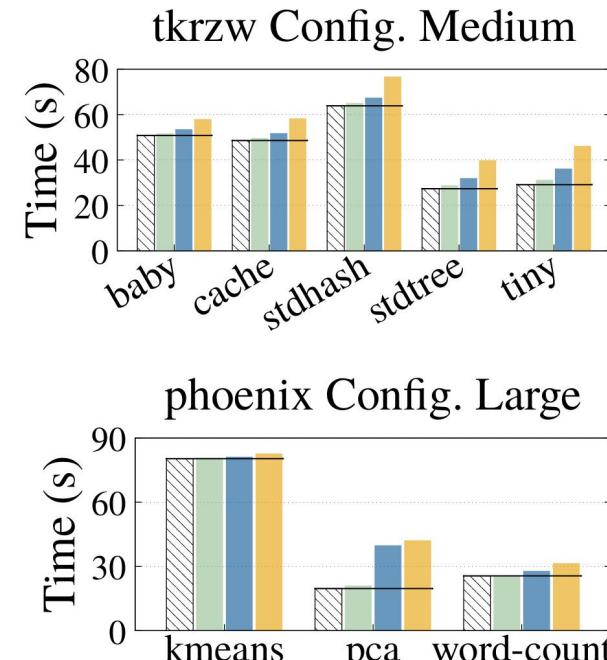
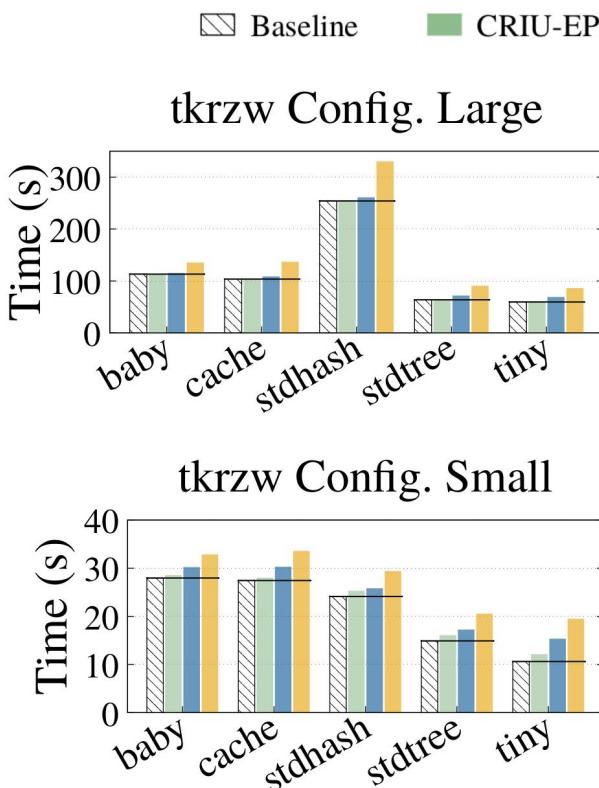
EPML:

- **13X** speedup vs. SPML for tkrzw-baby (config. Large)
- **4X** speedup vs. /proc for tkrzw-baby (config. Large)

Efficient Dirty Page Tracking In Virtualized Clouds



Evaluation Results with CRIU: for Tracked



CRIU pauses applications during checkpointing resulting in:

SPML:

- Up to **102%** overhead on Phoenix-pca with CRIU
- **34.6% overhead on average**

/proc:

- Up to **114%** overhead on Phoenix-pca with CRIU
- **15.6% overhead on average**

EPML:

- Only **7% overhead** on Phoenix-pca with CRIU
- **3.5% overhead on average**
- avg. **+9x** and **+4x improvement** on SPML & /proc resp.

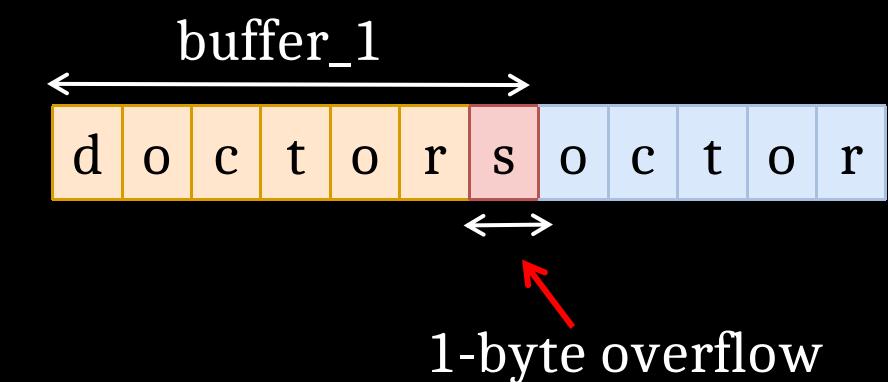
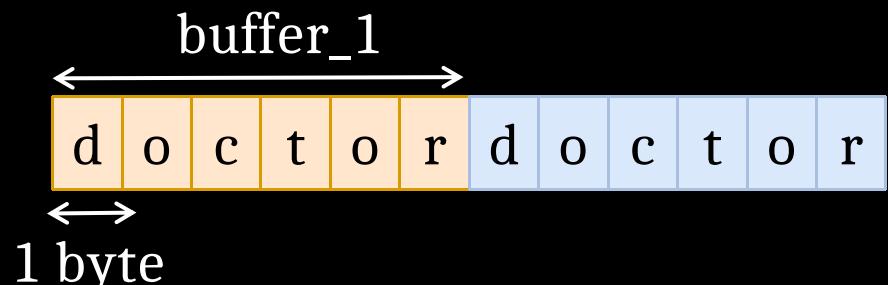
OoH for SPP: Efficient Buffer Overflow Detection In Virtualized Clouds

Efficient Buffer Overflow Detection In Virtualized Clouds



Buffer Overflow: Introduction

Definition

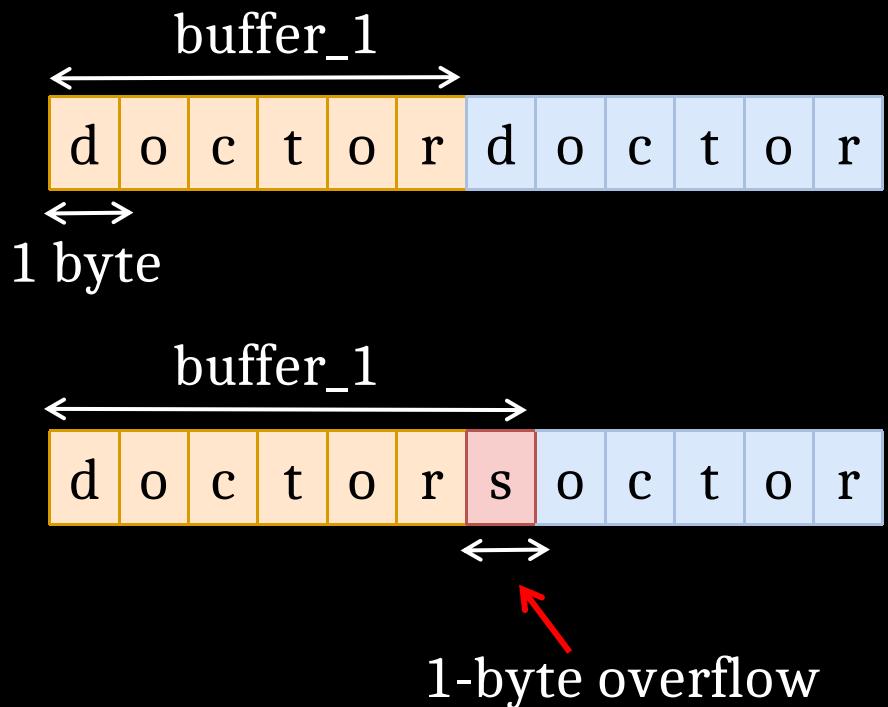


Efficient Buffer Overflow Detection In Virtualized Clouds



Buffer Overflow: Introduction

Definition



Importance

- ✓ Can lead, for example, to DB corruption (e.g., if overflow modifies a password), carrying out an attack (e.g., malware code injection), etc.
- ✓ 70% Google Chrome's bugs [50, Google security blog]
- ✓ **Most prevalent vulnerability in 2022 -- sources** [51, CWE (Common Weaknesses and Exposure)]

Efficient Buffer Overflow Detection In Virtualized Clouds



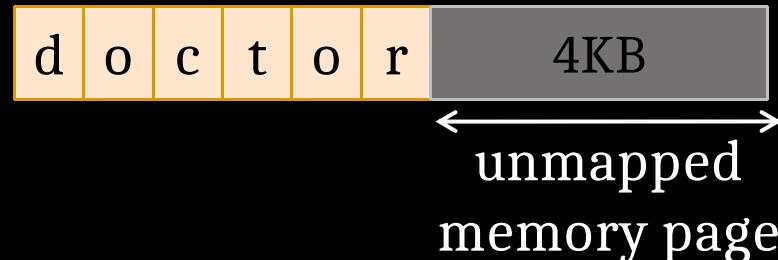
Buffer Overflow: State-of-the-art Mitigation Mechanisms

Detection Mechanisms

Canary



Guard Page



Hardware
Capabilities

CHERI: Capabilities Hardware Enhance RISC-V Instructions [159, ARCH 2014]

- Overflow detection by the HW
- Needs to rebuild systems [156, Tech. Rep. 2020]
- Incurs significant overheads (up to 250%) [156, Tech. Rep. 2020]

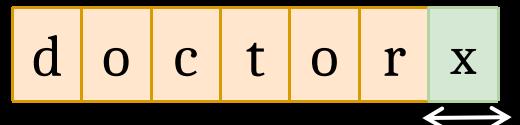
Efficient Buffer Overflow Detection In Virtualized Clouds



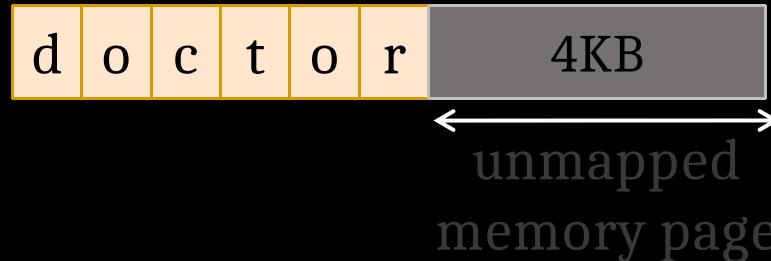
Buffer Overflow: State-of-the-art Mitigation Mechanisms

Detection Mechanisms

Canary



Guard Page



Hardware
Capabilities

CHERI: Capabilities Hardware Enhance RISC-V Instructions [159, ARCH 2014]

- Overflow detection by the HW
- Needs to rebuild systems [156, Tech. Rep. 2020]
- Incurs significant overheads (up to 250%) [156, Tech. Rep. 2020]

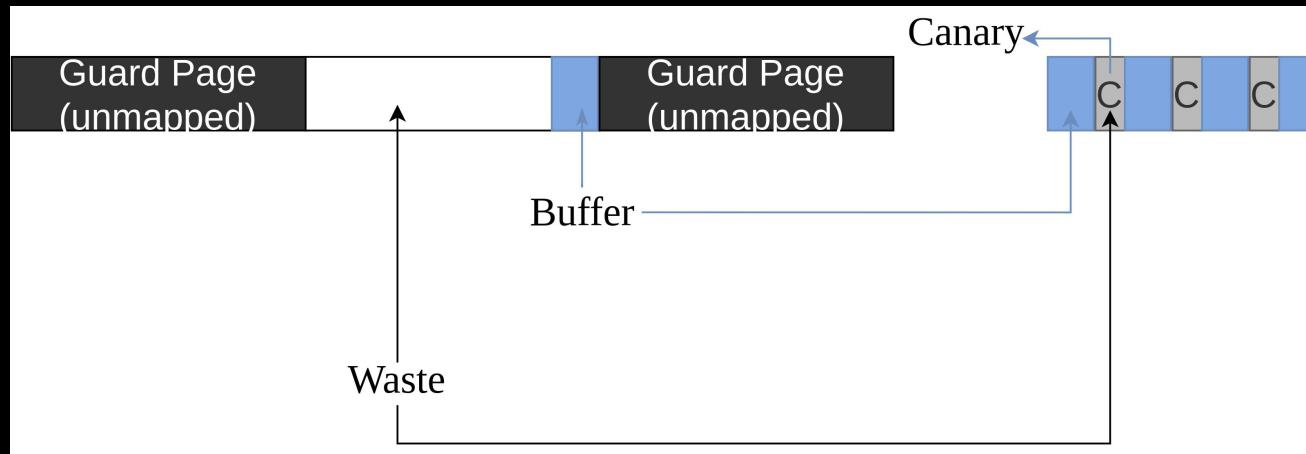
Efficient Buffer Overflow Detection In Virtualized Clouds



Synchronous Detection vs Memory Overhead

Canary
modest memory overhead
asynchronous detection

Guard page
significant memory overhead
synchronous detection



Efficient Buffer Overflow Detection In Virtualized Clouds

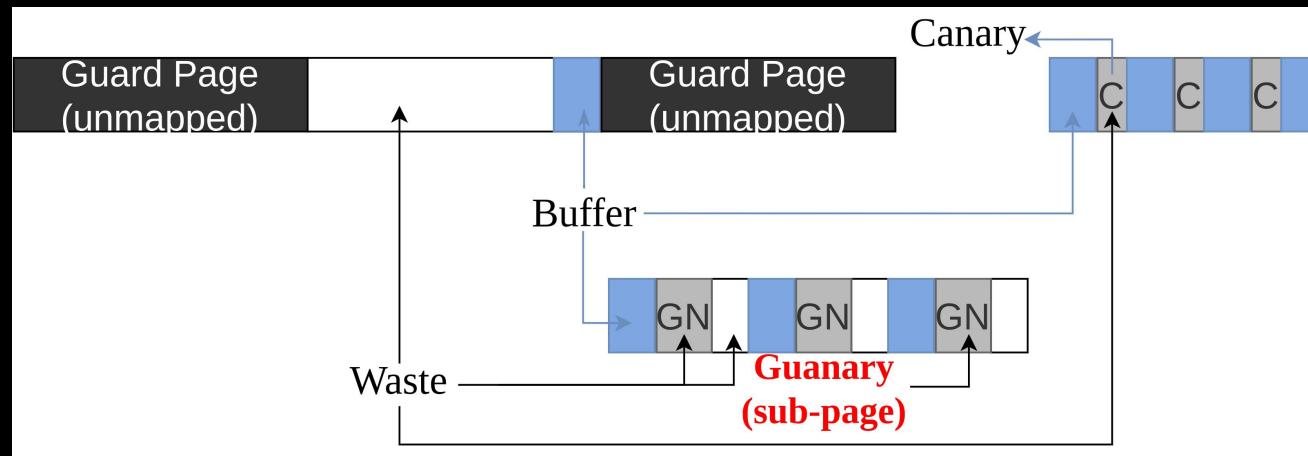


Synchronous Detection vs Memory Overhead

Canary
modest memory overhead
asynchronous detection

GuaNary
modest memory overhead
synchronous overflow detection

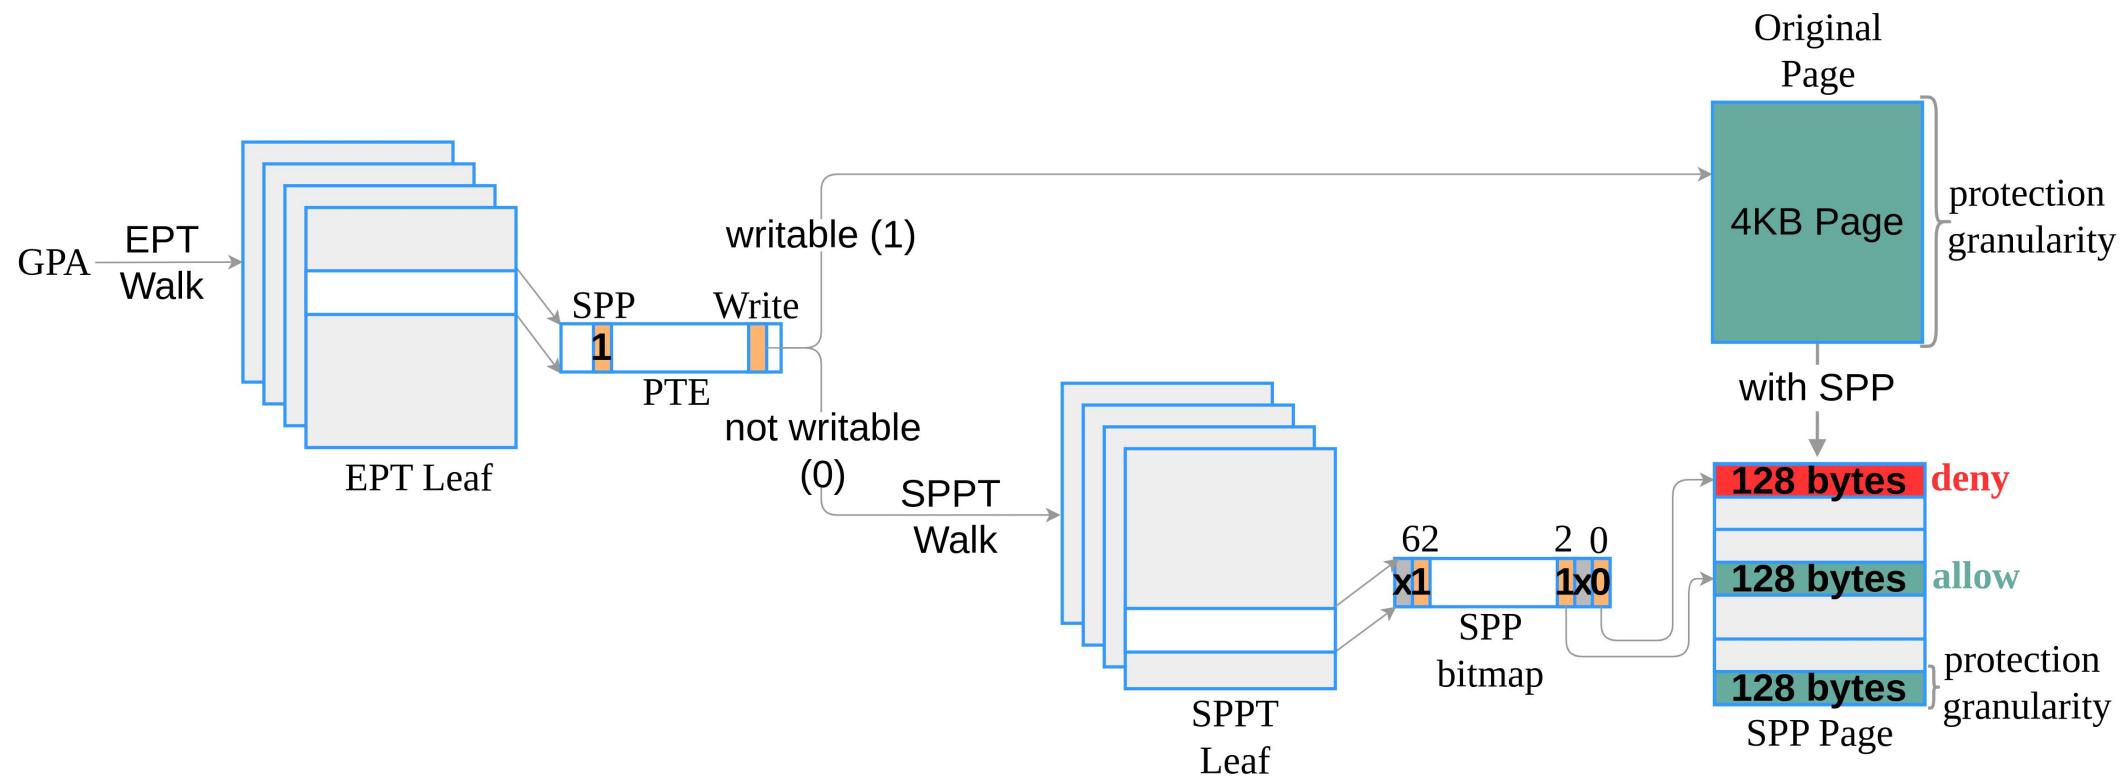
Guard page
significant memory overhead
synchronous detection



Efficient Buffer Overflow Detection In Virtualized Clouds



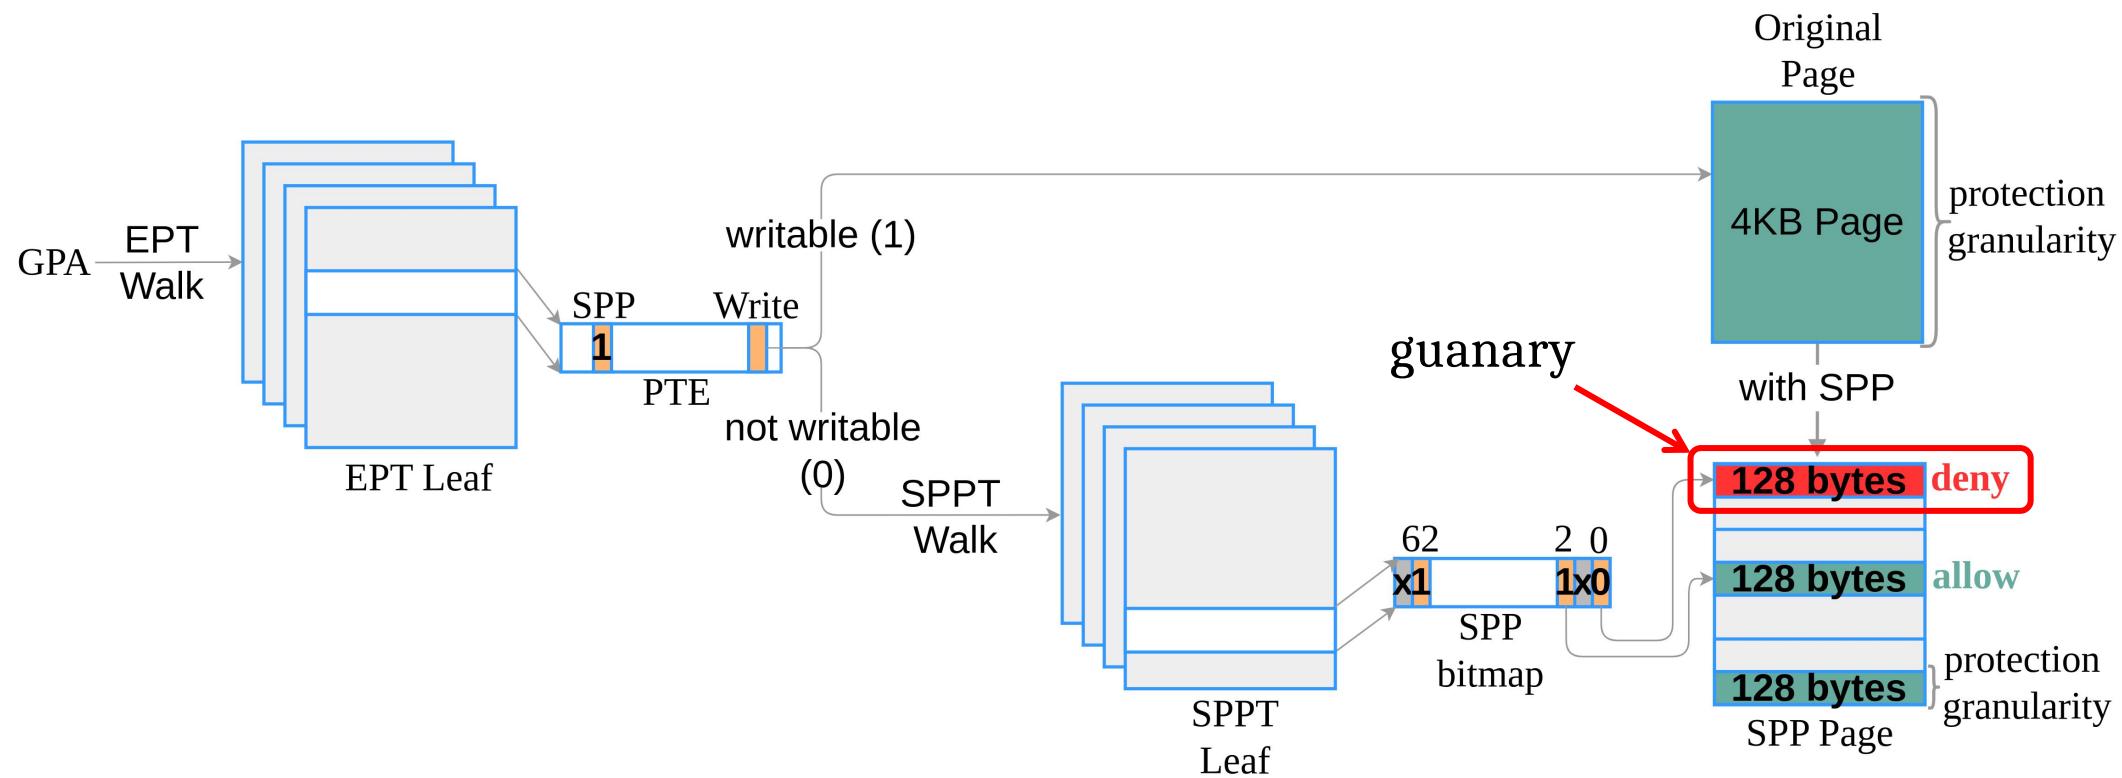
GuaNary: Intel Sub-Page write Permissions (SPP)



Efficient Buffer Overflow Detection In Virtualized Clouds



GuaNary: Intel Sub-Page write Permissions (SPP)



Efficient Buffer Overflow Detection In Virtualized Clouds



GuaNary: Challenges

Can be summarized in 2 main:

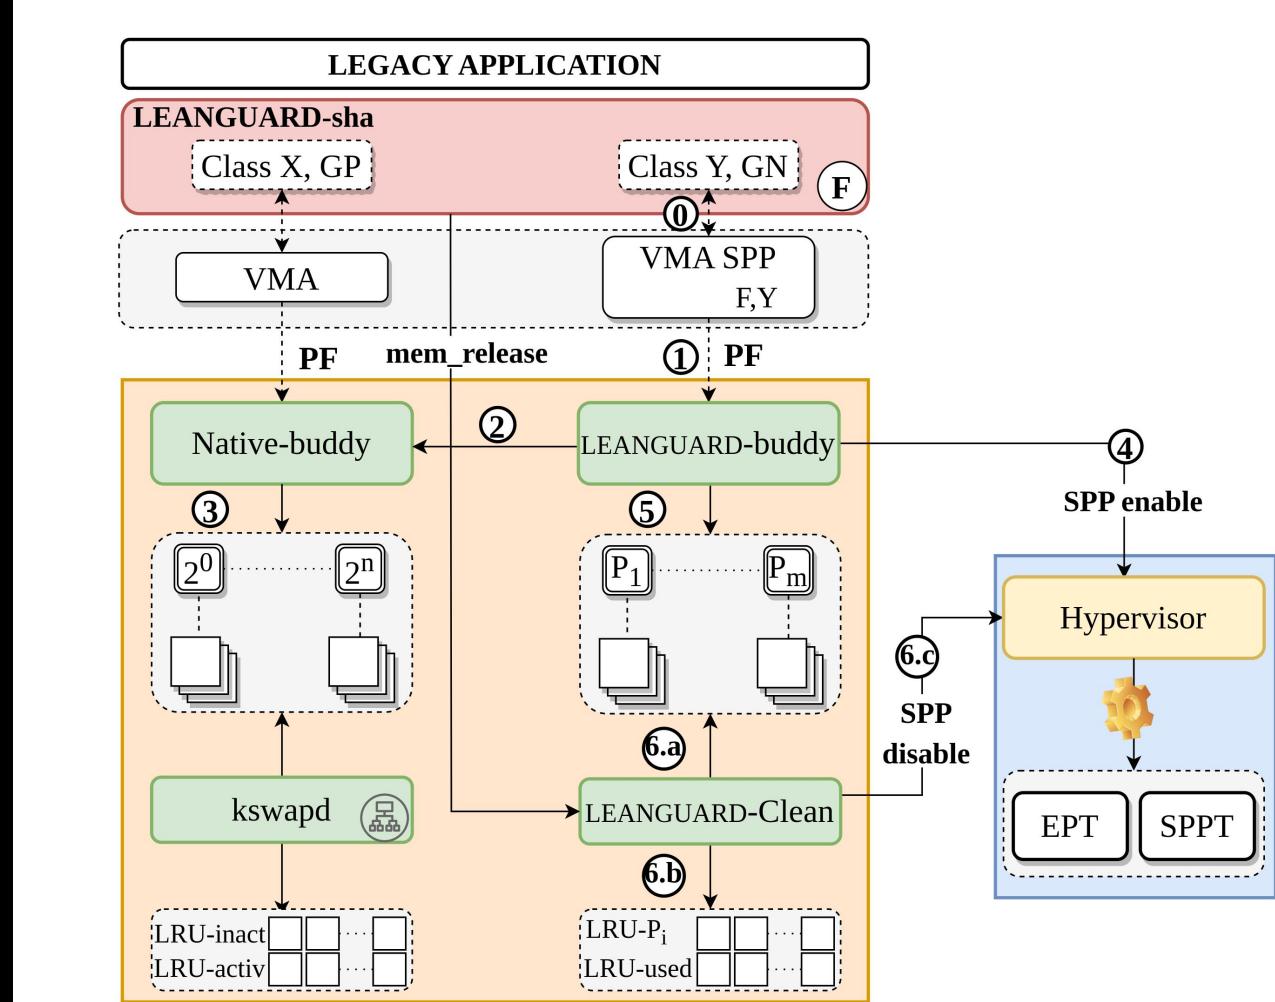
Conceptual: Costly hypercalls and EPT+SPPT walk

Technical: Page heterogeneity

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard: Complete SW Stack to use Guanary



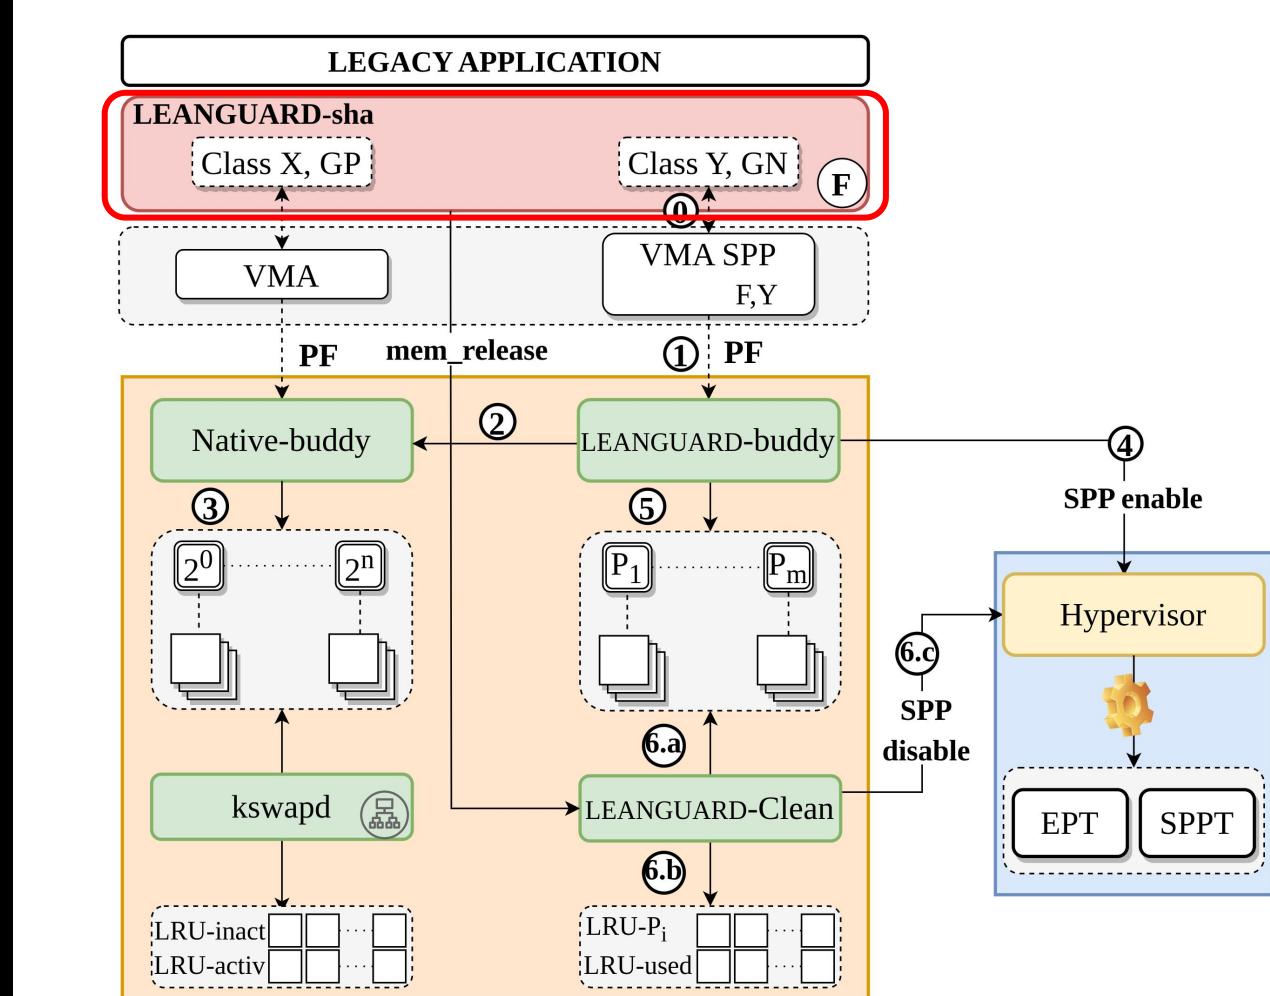
Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard: Complete SW Stack to use Guanary

LeanGuard-sha configures 2 param:

- G: type of guardian (GP or GN)
 - GP if $\text{sizeof(class)} \% 4\text{KB} = 0$
 - Guanary if not
- F: protection frequency (configured by the user)

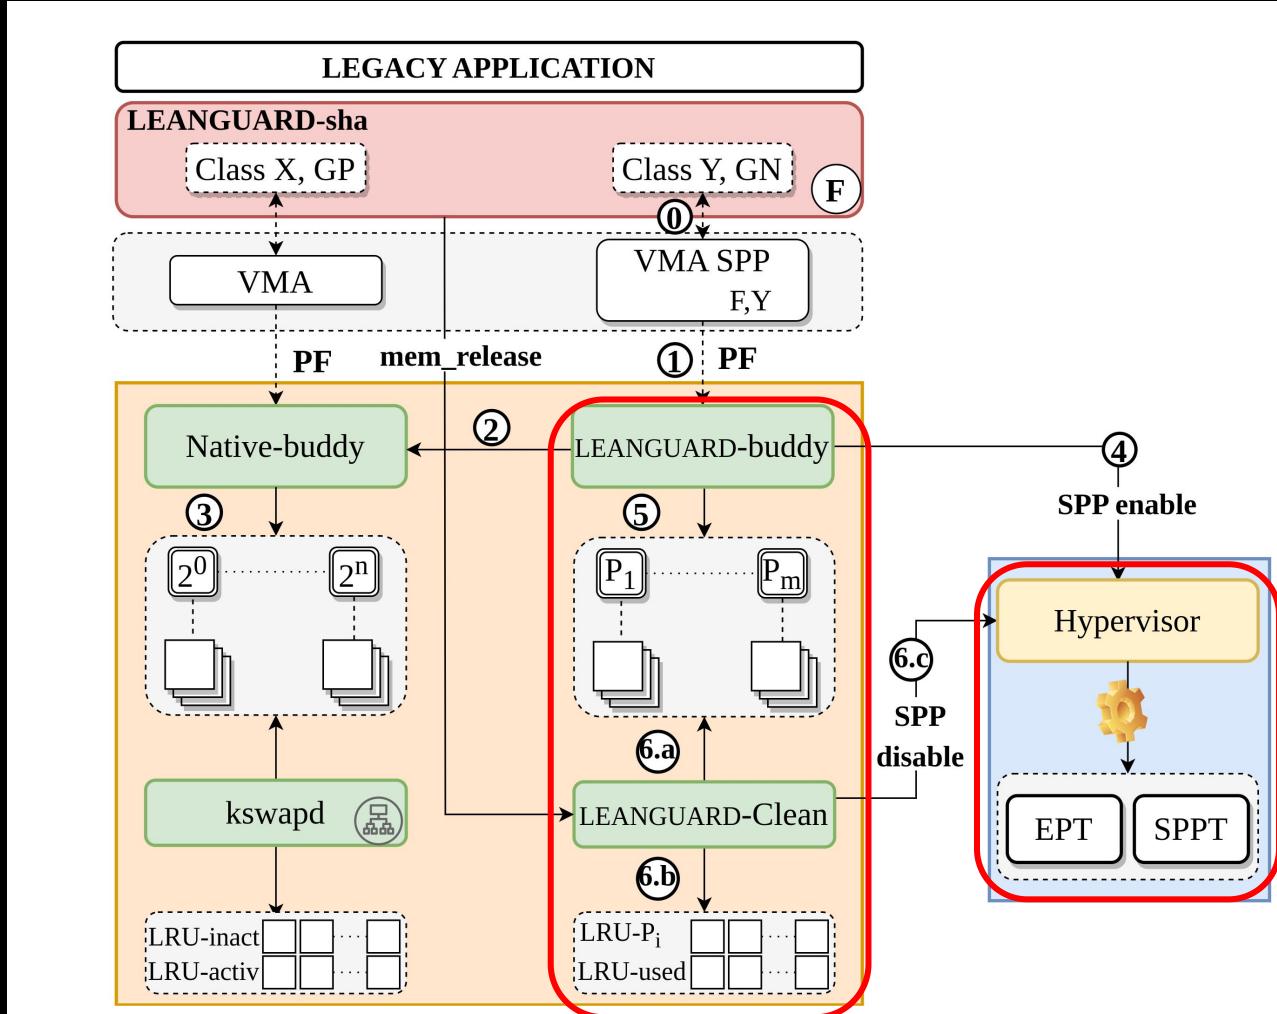


Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard: Complete SW Stack to use Guanary

- New buddy allocator responsible for SPP page allocation
- Batching hypercalls for optimal performance
- Hypervisor optimization: SPP config for a group of contiguous pages



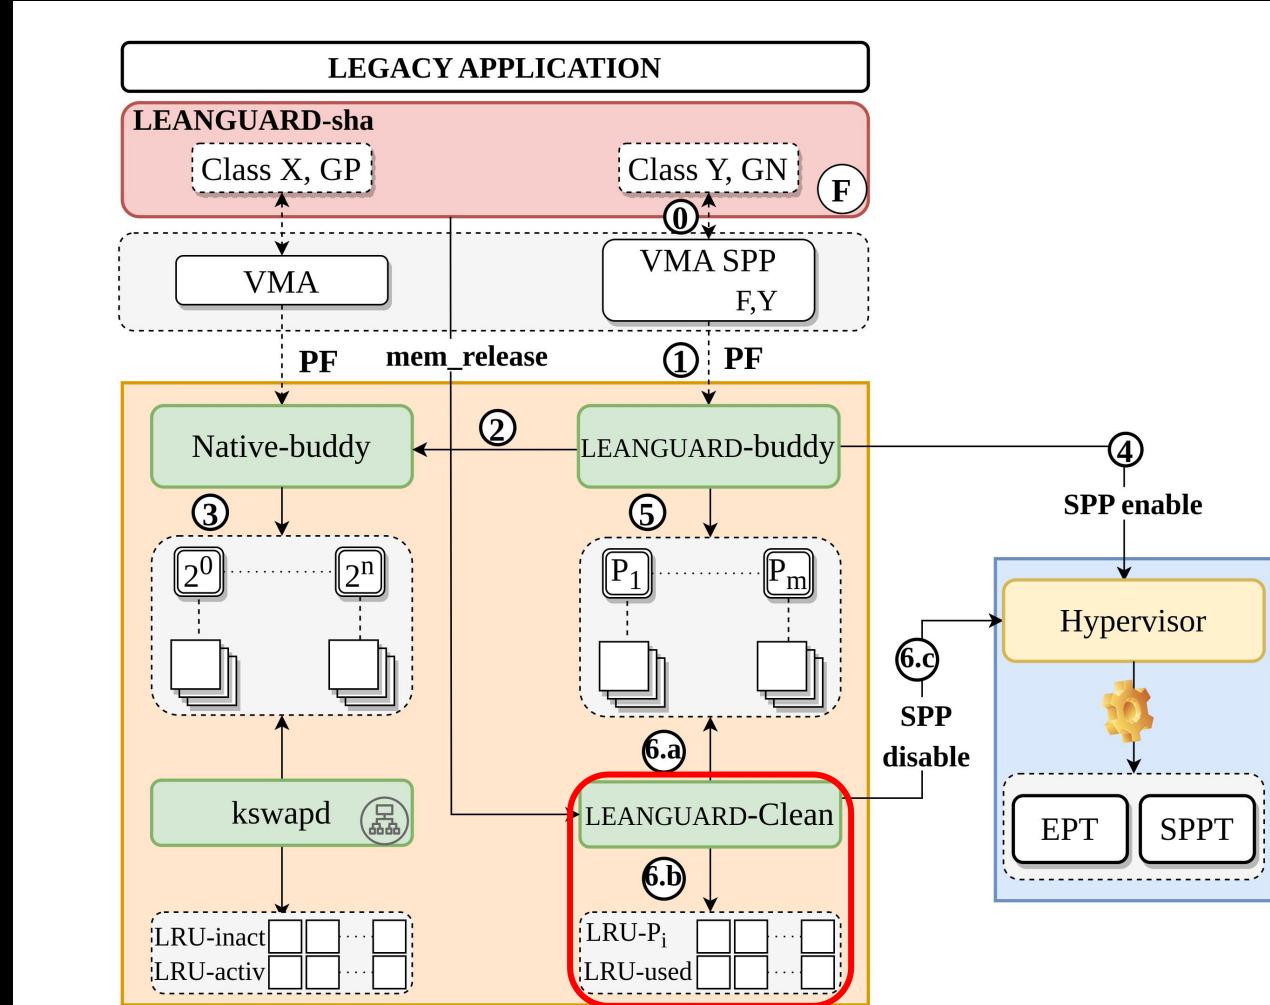
Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard: Complete SW Stack to use Guanary

Memory reclaim: LeanGuard-Cleaner

- Handles SPP page release on process termination
 - Reinserts SPP pages into corresponding pool
 - Returns the excess of pages to Native-buddy for all pools: batch hypercall for SPP deactivation



Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard: Implementation

Xen hypervisor version 4.10: 600LOCs addition

Linux OS version 5.11.14: 750LOCs addition

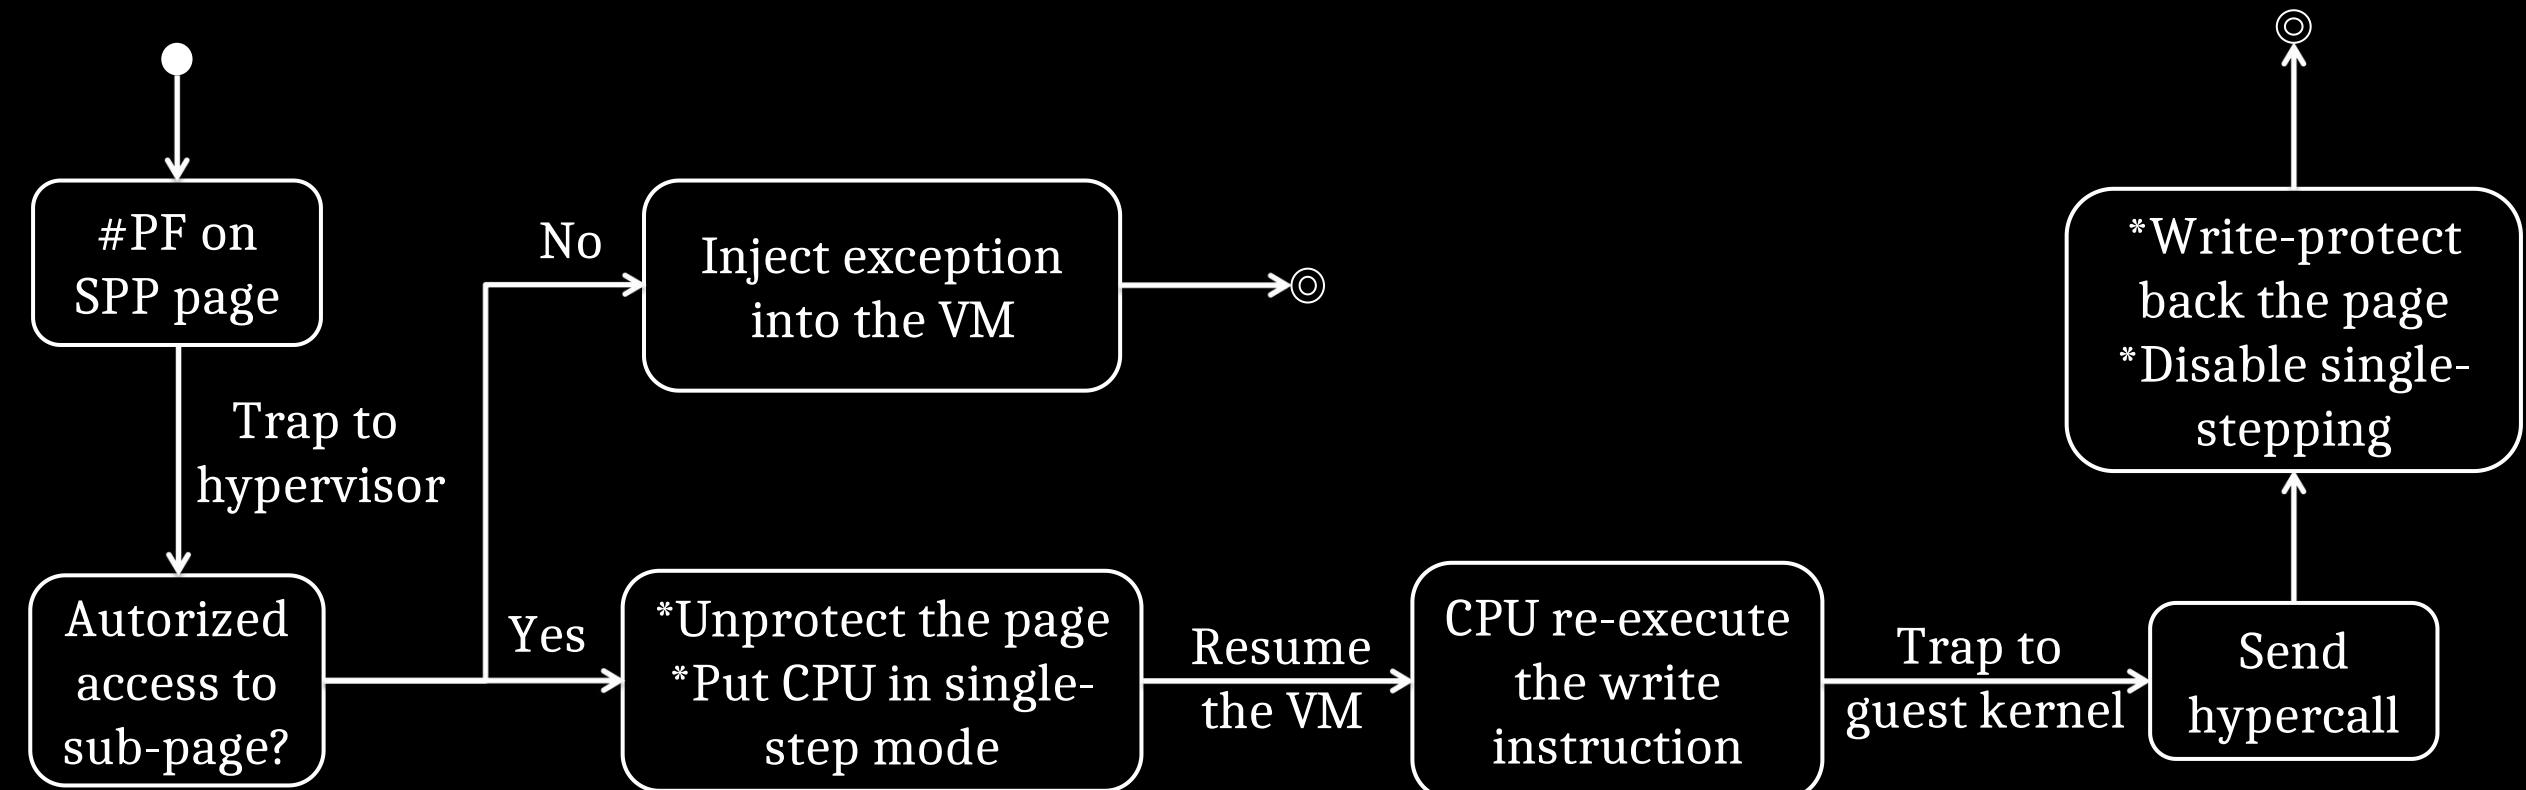
SlimGuard secure allocator: 100LOCs addition

- 2019 State-of-the-art BIBOP (Big Bag Of Pages) secure allocator
- Proven more memory-efficient than other state-of-the-art BIBOP allocators (Guarder, FreeGuard, etc.)

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Memory Overhead Methodology



Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation

Benchmarks

PARSEC applications: blackscholes, bodytrack, dedup, fluidanimate, freqmine, raytrace, streamcluster, swaptions, and x264

Our experiments hinge on:

Memory overhead and Security trade-off

Performance overhead and scalability

LeanGuard's goal is minimizing memory overhead while allowing synchronous buffer overflow detection

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Memory Overhead Results



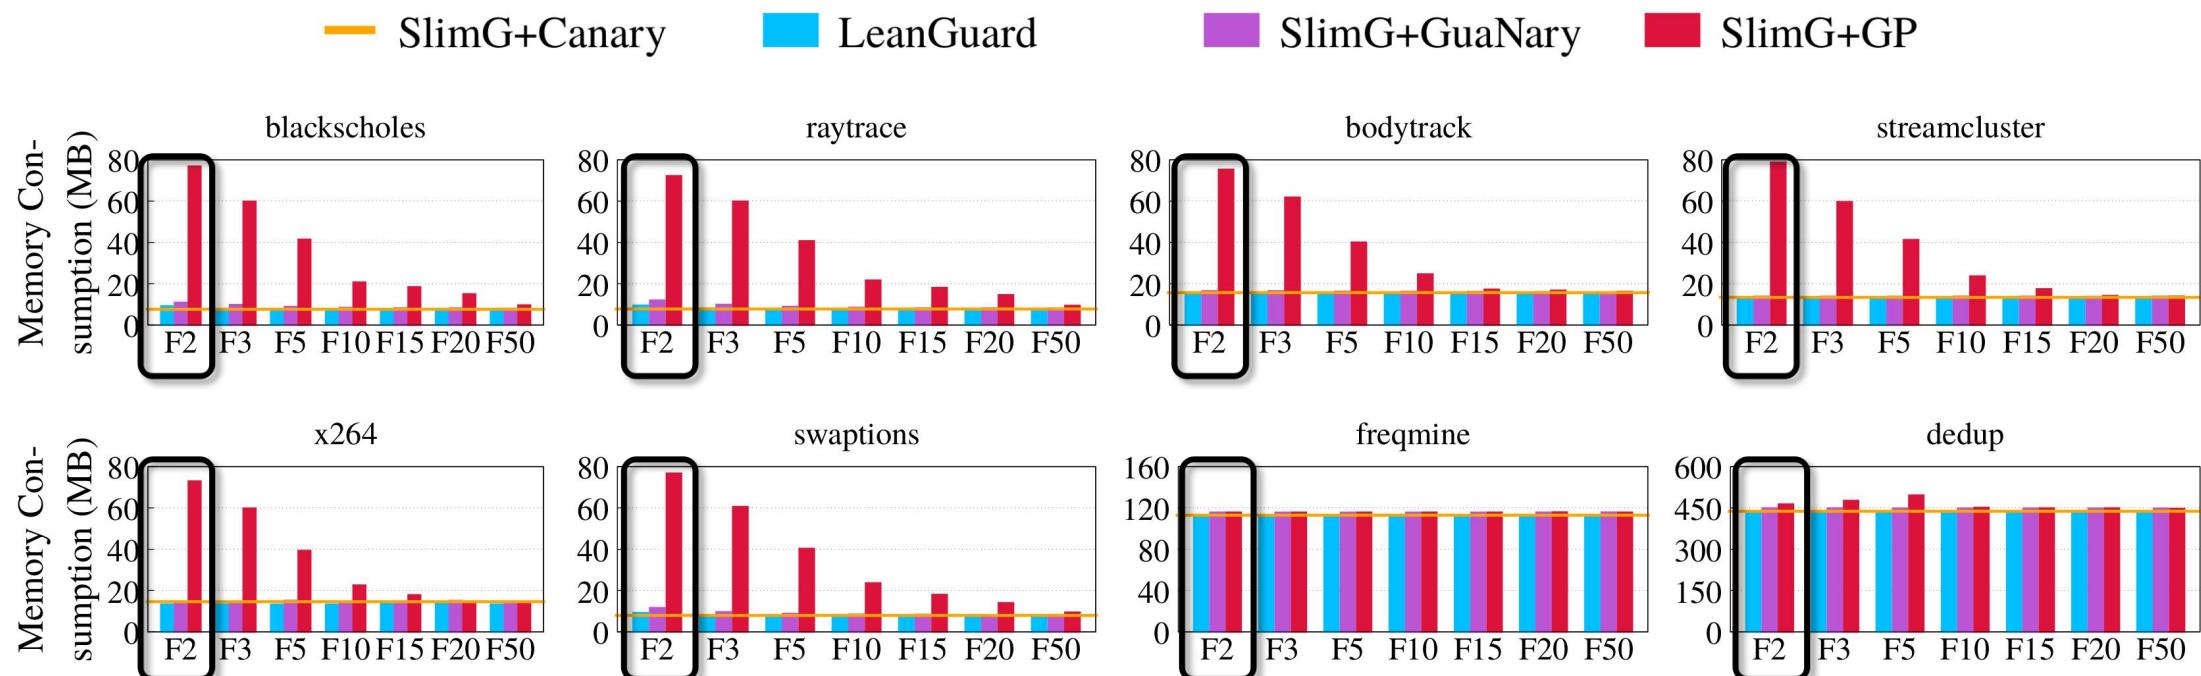
GuaNary effectively leads to memory consumption reduction

Slight improvement that LeanGuard brings to SlimG+GuaNary (26% swaptions, 25% on raytrace, and 18% blacksholes, when F=2) is explained by the fact that, unlike SlimG+GuaNary that exclusively uses GuaNary, LeanGuard can use 4KB guard pages when necessary

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Memory Overhead Results

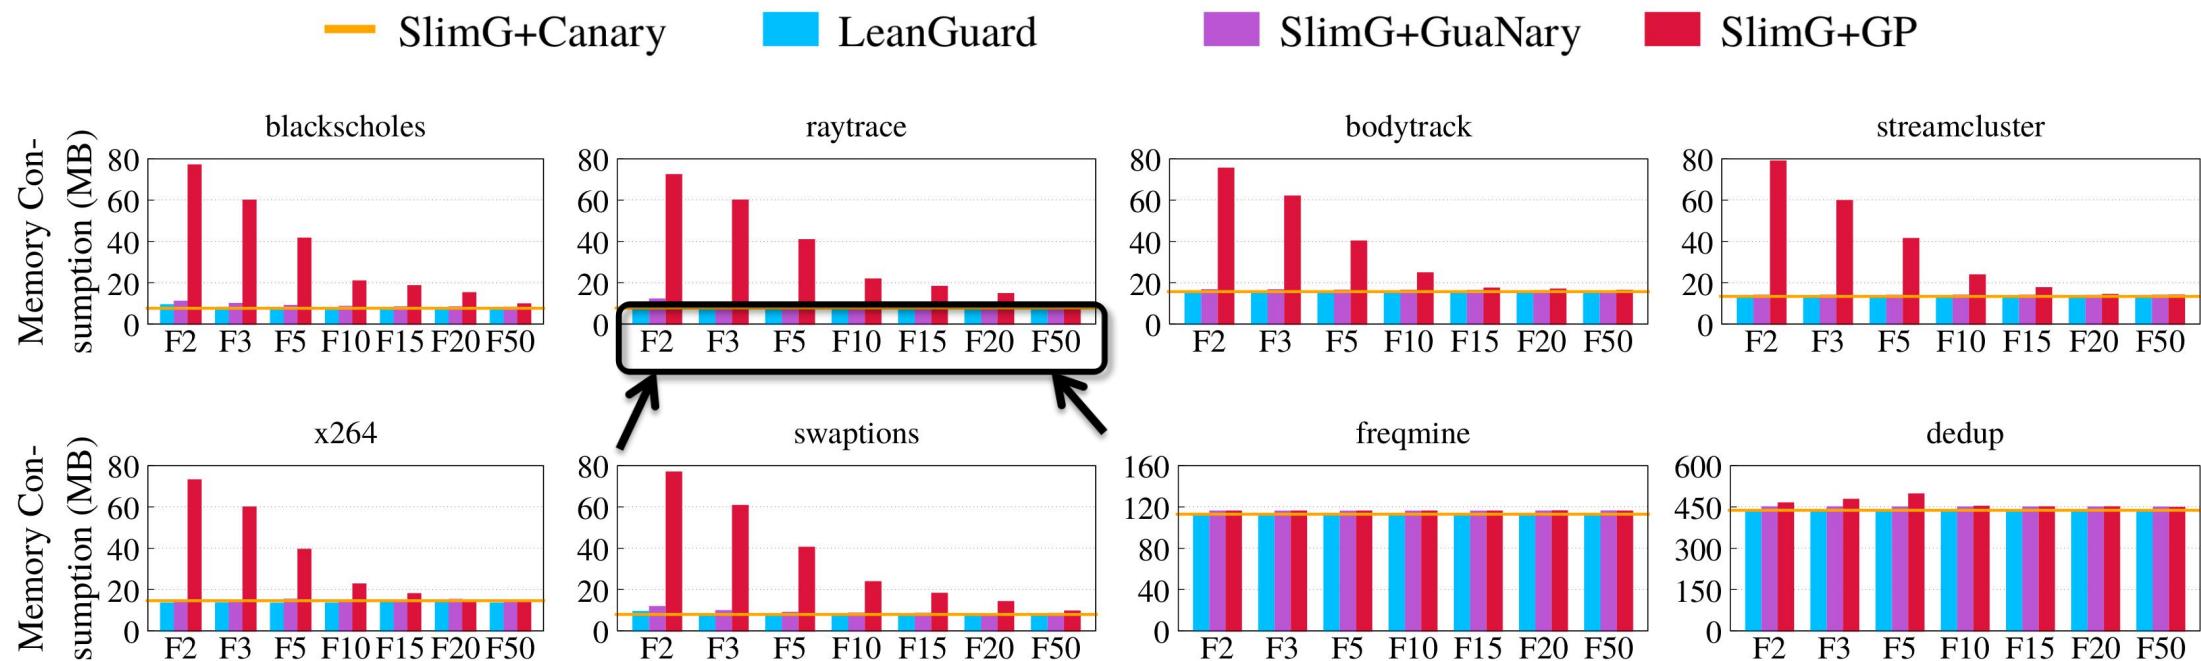


SlimG+GP incurs a significant memory overhead compared to LeanGuard. For example, to protect 50% of the allocated buffers (F=2), LeanGuard, on average, uses 60% less memory than SlimG+GP

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Memory Overhead Results



Using the same amount of memory as SlimGuard, LeanGuard allows protecting 25× more buffers than SlimG+GP

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Methodology

$$T_{est} = T_{noSPP} + \underbrace{TLB_{SPP} \times 4 \times T_{mem_latency}}_{\text{Capture the HW-part overhead}}$$



Capture the
SW overhead

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Methodology

$$T_{est} = T_{noSPP} + TLB_{SPP} \times 4 \times T_{mem_latency}$$



#TLB_misses due to SPP pages

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Methodology

$$T_{est} = T_{noSPP} + TLB_{SPP} \times 4 \times T_{mem_latency}$$

SPPT walk
requires 4
memory accesses

Time of memory
access ($\sim 111\text{ns}$)

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Methodology

$$T_{est} = T_{noSPP} + TLB_{SPP} \times 4 \times T_{mem_latency}$$



$$= k \times TLB_{tot}$$



total number of #PFs

obtained using the Linux' `perf` tool

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Methodology

$$T_{est} = T_{noSPP} + TLB_{SPP} \times 4 \times T_{mem_latency}$$



$$= k \times TLB_{tot}$$



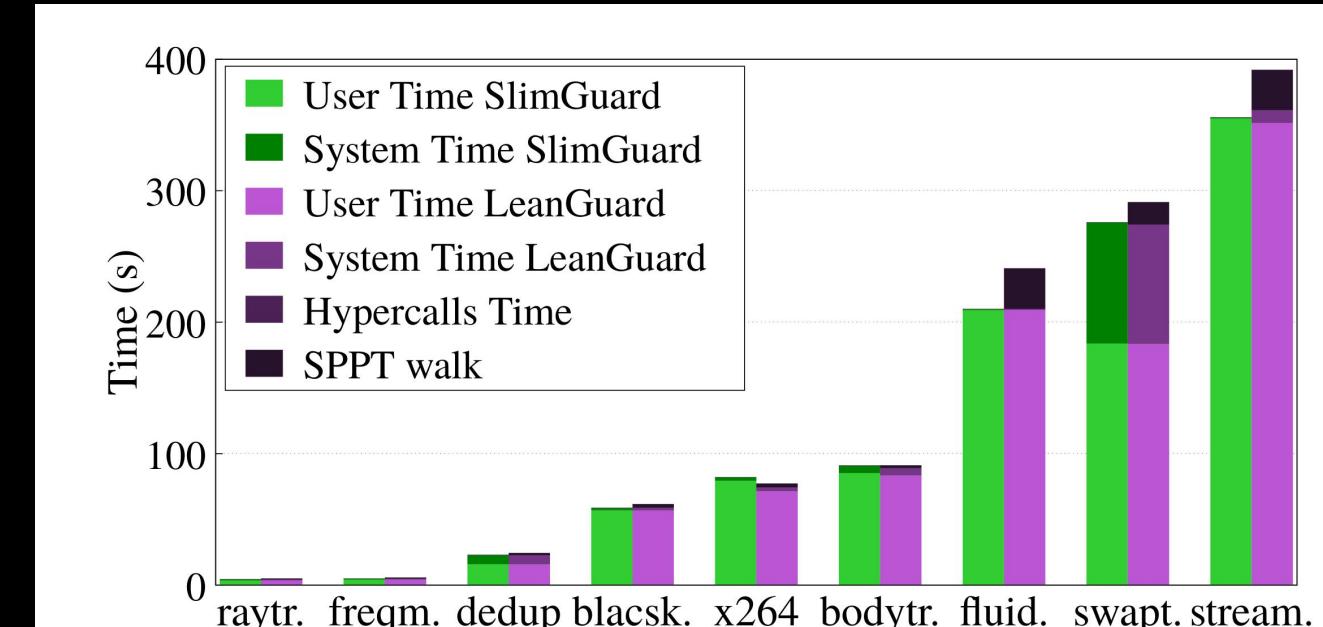
% of SPP #PFs

obtained by instrumenting Leanguard-buddy and Native-buddy

Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Results



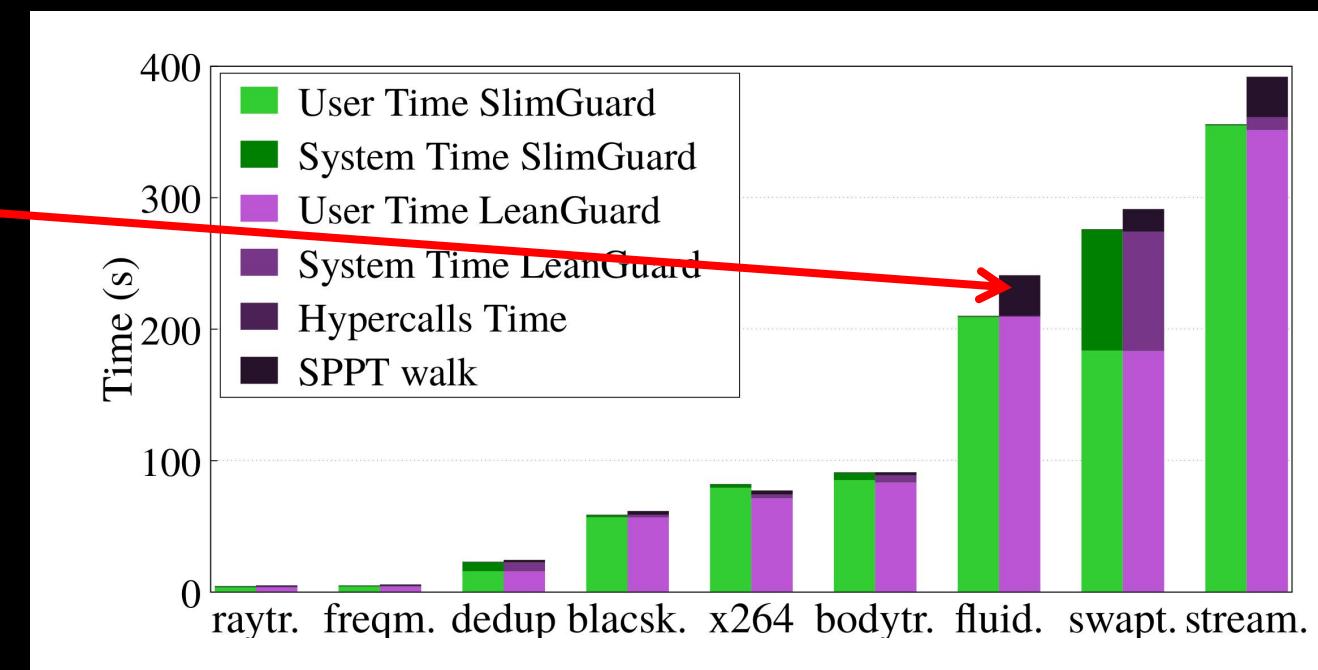
Efficient Buffer Overflow Detection In Virtualized Clouds



LeanGuard Evaluation: Performance Overhead Results

The gap between SlimGuard and LeanGuard is essentially the SPPT walk

LeanGuard incurs only ~7.7% overhead on average: can be acceptable to a user whose priority is security



Summary and Future Directions

Summary

1. Introduction to OoH (Out of Hypervisor):
 - New virtualization research axis
 - Exposure of hypervisor-oriented HW virt. features to guest processes
 - Excellent alternative to nested virtualization
2. The large ecosystem of hardware virtualization features can help improving existing software-based tools
3. Think of OoH and virt. feature categorization from the conception / design of HW virtualization features

Future Directions

OoH in bare-metal:

- ✓ Make host userspace applications take advantage of HW virtualization functionalities

OoH for other hypervisors, processors, and architectures:

- ✓ OoH' principles can be applied to processors other than Intel (e.g., AMD)
- ✓ OoH can find applicability to AArch64 architectures (e.g., with NEVE)



Publications and Awards

Publications		Awards
International	National	
Out of Hypervisor (OoH): Efficient Dirty Page Tracking In Userspace Using Hardware Virtualization Features, S. Bitchebe & A. Tchana, SC 2022	Applying Intel SPP Feature for Buffer Overflows Mitigation, S. Bitchebe et al., Compas 2022	2022 Google Scholarship for Women in Computer Science
Extending Intel PML for Hardware-Assisted Working Set Size Estimation of VMs, S. Bitchebe et al., VEE 2021	Caracterization and Methodology of Buffer Overflow Analysis, Y. KONE, S. Bitchebe & A. Tchana, Compas 2022	2021 Microsoft Research Ph.D. Fellowship
Study of Intel PML Effectiveness, S. Bitchebe et al., Eurosyst Doctoral Workshop 2019	sQemu: Toward a Scalable Virtual Storage, K. Nguetchouang, S. Bitchebe et al., Compas 2022	2021 L' Oréal-UNESCO For Women in Science Program
Rejects Eurosyst 2023, MICRO 2019, JSS 2020, Eurosyst 2020, SIGMETRICS 2020, IFIP PERF 2020, ISCA 2022	Hardware Assisted Virtual Machine Page Tracking, S. Bitchebe et al., Compas 2019	2021 NEC Lab Ph.D. Research Fellowship